



POLITECNICO
MILANO 1863

SYSTEMS AND METHODS FOR BIG AND UNSTRUCTURED DATA

Towards NoSQL

Marco Brambilla

marco.brambilla@polimi.it







 @marcobrambi

Agenda

- Schema
- Scalability
- Transactional Properties
- CAP Theorem
- ACID vs. BASE
- NoSQL models and solutions

Towards Big Data Architectures

Big Data vs. Traditional Data

	Traditional	Big Data
Data Characteristics	 Relational (with highly modeled schema)	
Cost	 Expensive (storage and compute capacity)	 Commodity (storage and compute capacity)
Culture	 Rear-view reporting (using relational algebra)	 Intelligent action (using relational algebra AND ML, graph, streaming, image processing)

One keyword:

FLEXIBILITY


Schema

The schema-less idea

- Schema-less or implicit schema?
 - Aggregate-based
 - Key-value and document-based
 - Relationship-based
 - Graph DBs are better than relational!
- Ref. Domain-Driven Design, by Eric Evans.


Paradigmatic shift introduced by Big Data: from Schema On Write ...

- Long lasting discussion about the schema that can accommodate all needs
- Some analysis can no longer be performed because the data were lost at writing time

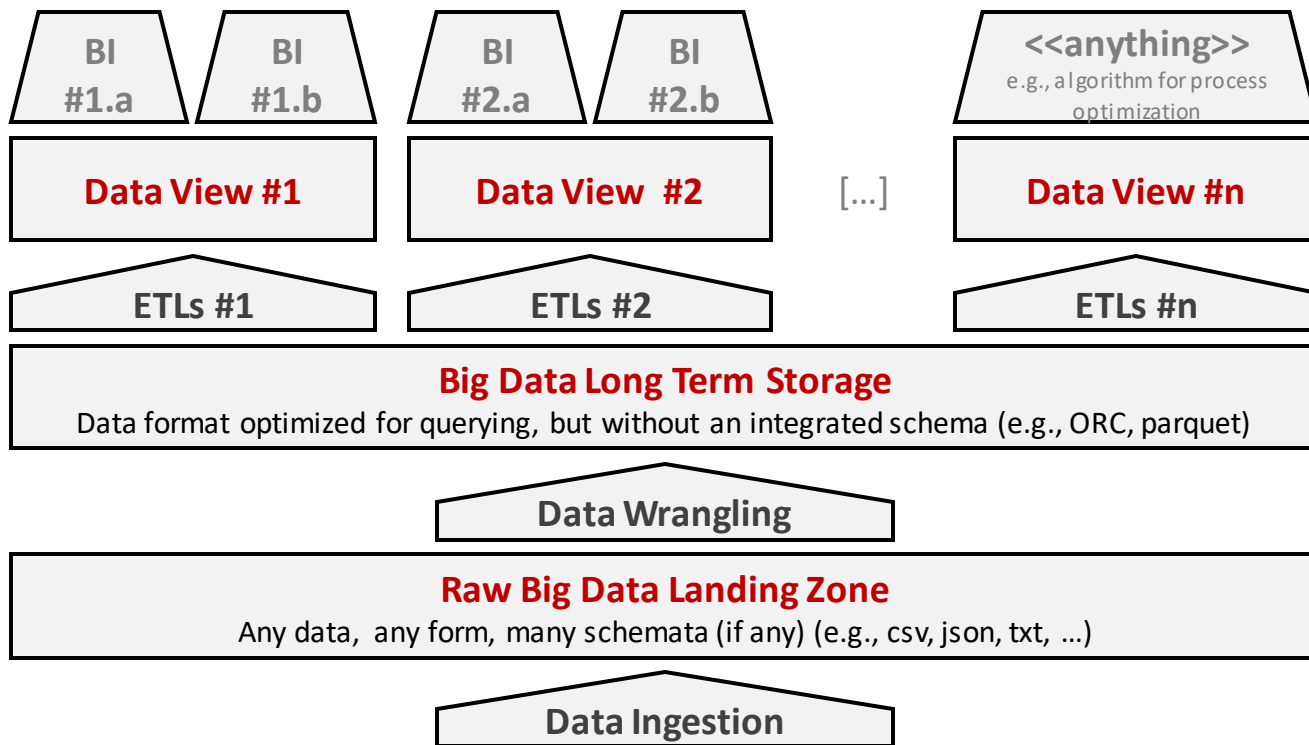


Paradigmatic shift introduced by Big Data: ... to Schema On Read

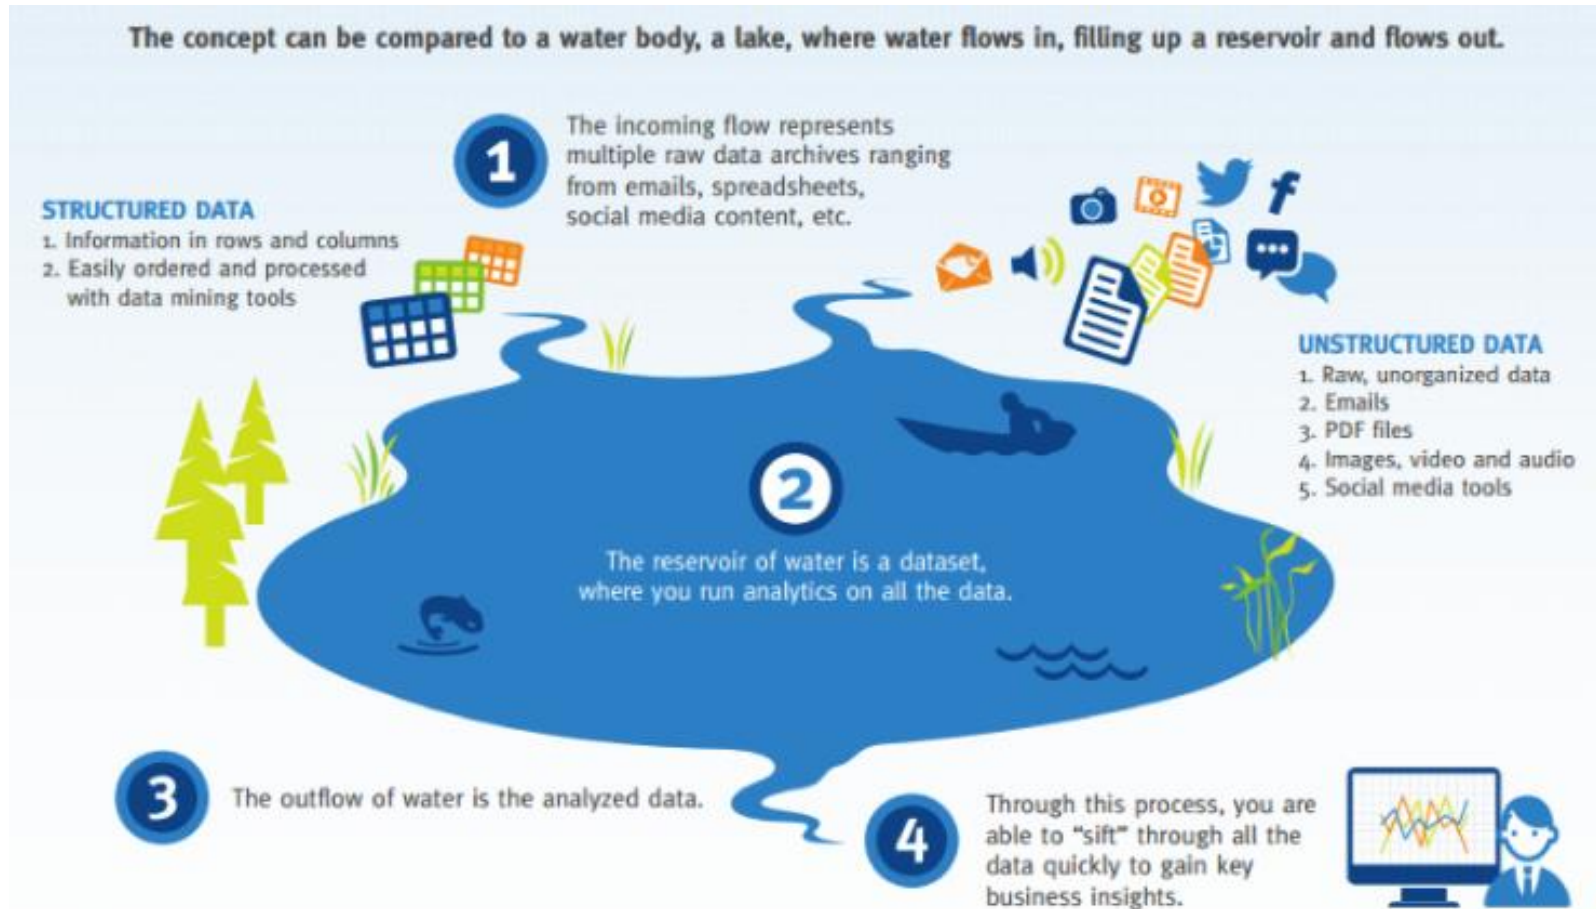
- Load data first, ask question later
- All data are kept, the minimal schema needed for an analysis is applied when needed
- New analyses can be introduced in any point in time



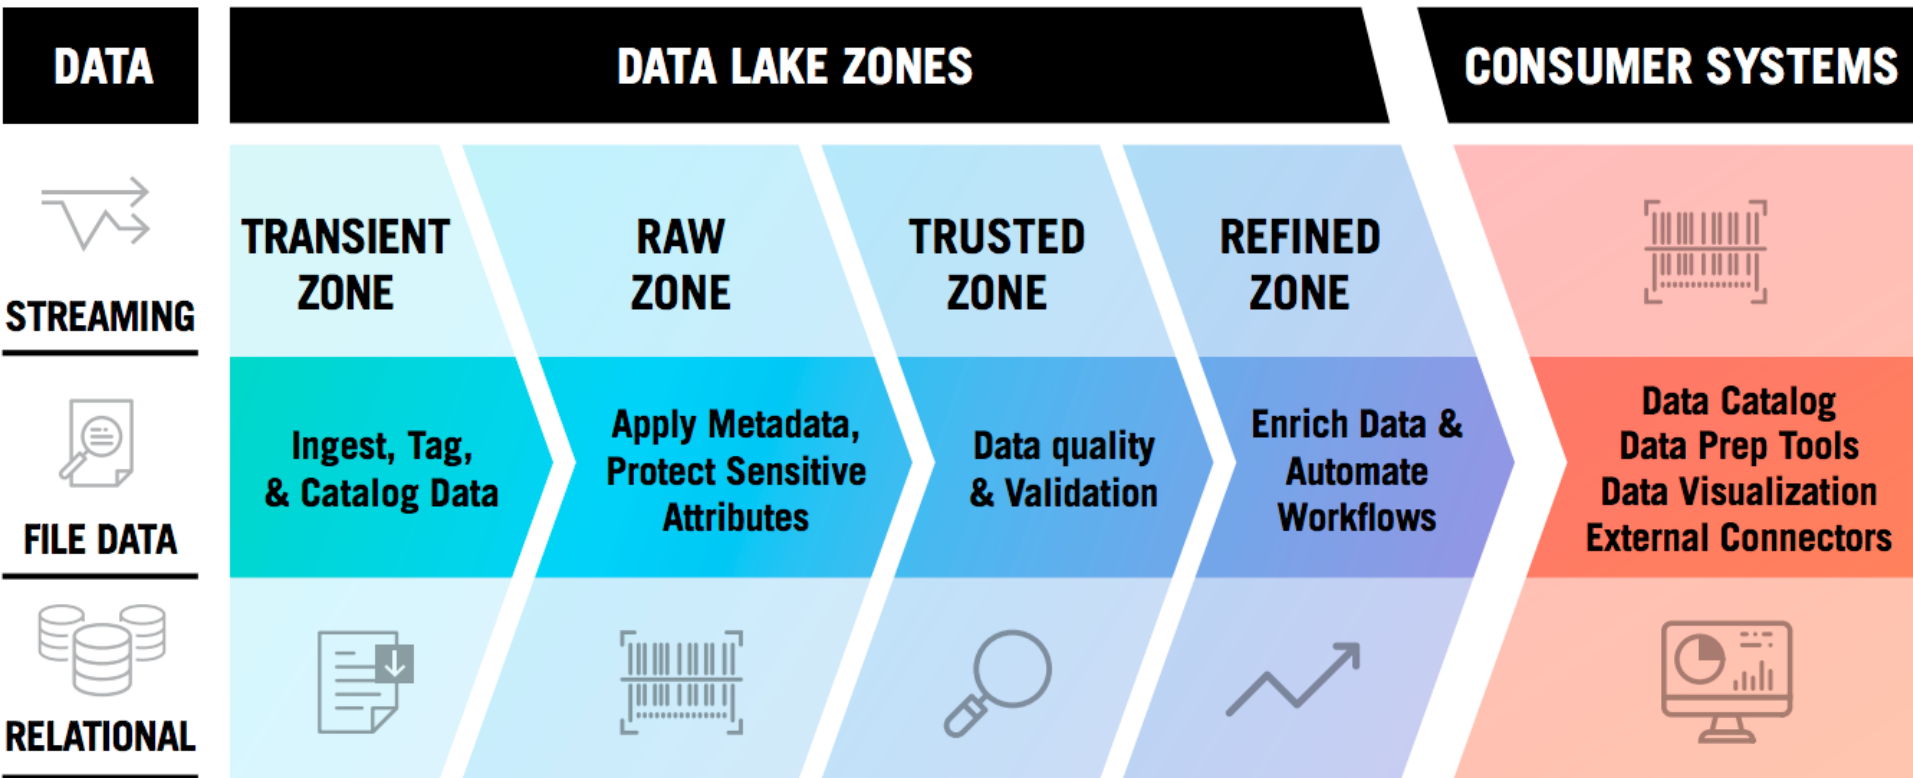
Logical Architecture of Big Data



The Concept of Data Lake



Data Lakes in Process




Another issue: ORM

- Impedance Mismatch
- Object-Relational Mapping problem (and solution)
- Object Orientation (OODB)
 - Commercial failure!

Extract Transform Load (in the Big Data era)

The process that extracts data from heterogeneous data sources, transforms it in the schema that better fits the analysis to perform and loads it in the system that will perform the analysis.




Scalability for Big Data

Vertical Vs. Horizontal Scalability

- "Traditional" SQL system scale vertically:
 - Adding data to a "traditional" SQL system may degrade its performances
 - When the machine, where the SQL system runs, no longer performs as required, the solution is to buy a better machine (with more RAM, more cores and more disk)
- Big Data solutions scale horizontally
 - Adding data to a Big Data solution may degrade its performances
 - When the machines, where the big data solution runs, no longer performs as required, the solution is to add another machine

Vertical scalability




Commodity hardware

- CPU: 8-32 cores
- RAM: 16-64 GB
- Disk: 1-3 TB
- Network: 10 GE




Appliance


- e.g. ORACLE EXADATA DATABASE MACHINE X6-8
- CPU: 576 cores
- RAM: 24TB
- Disk: 360TB of Flash Storage per rack
- Network: 40 Gb/second InfiniBand




Horizontal scalability




Vertical vs. Horizontal scalability



Vertical vs. Horizontal scalability



The "grey area" is time dependent




Transactional Properties

In OLTP systems

Definition of Transaction

- An elementary unit of work performed by an application
- Each transaction is encapsulated within two commands:
 - **begin transaction** (bot) and **end transaction** (eot)
- Within a transaction one of the commands below is executed (exactly once):
 - **commit work** (commit) and **rollback work** (abort)
- **Transactional System** (OLTP): a system capable of providing the definition and execution of transactions on behalf of multiple, concurrent applications

Application and Transaction




ACID Properties of Transactions

- A transaction is a unit of work enjoying the following properties:
 - Atomicity
 - Consistency
 - Isolation
 - Durability

Atomicity

- A transaction is an atomic transformation from the initial state to the final state
- Possible behaviors:
 1. Commit work: SUCCESS
 2. Rollback work or error prior to commit: UNDO
 3. Fault after commit: REDO




Consistency

- The transaction satisfies the integrity constraints
- Consequence:
 - If the initial state is consistent
 - Then the final state is also consistent

Isolation

- A transaction is not affected by the behavior of other, concurrent transactions
- Consequence:
 - Its intermediate states are not exposed
 - The “domino effect” is avoided



Durability

- The effect of a transaction that has successfully committed will last “forever”
 - Independently of any system fault

Transaction Properties and Mechanisms

- Atomicity
 - Abort-rollback-restart
 - Commit protocols
- Consistency
 - Integrity checking of the DBMS
- Isolation
 - Concurrency control
- Durability
 - Recovery management

CAP Theorem

CAP Theorem (Brewer's Theorem)

It is impossible for a [distributed computer system](#) to simultaneously provide all three of the following guarantees:

- [Consistency](#): all nodes see the same data at the same time
 - [Availability](#): Node failures do not prevent other survivors from continuing to operate (a guarantee that every request receives a response about whether it succeeded or failed)
 - [Partition tolerance](#): the system continues to operate despite arbitrary partitioning due to network failures (e.g., message loss)
- A distributed system can satisfy any two of these guarantees at the same time but not all three.

CAP Theorem (Brewer's Theorem): So, what can be done?

In a distributed system, a network (of networks) is inevitable (by definition).

Failures can, and will, occur to a networked system -> **partitioned tolerance should be accommodated**.

Then, the only option left is **choosing between Consistency and Availability** - i.e., CA doesn't make any sense (except when we have, e.g., a single-site databases; 2-phase commit, cache validation protocols)

Not necessarily in a mutually exclusive manner, but possibly by partial accommodation of both
→ trade-off analysis important


CP: A partitioned node returns

- a correct value, if in a consistent state;


- a timeout error or an error, otherwise

AP: A partitioned node returns the most recent version of the data, which could be stale.

CAP: what's really about?



CAP: what's really about?



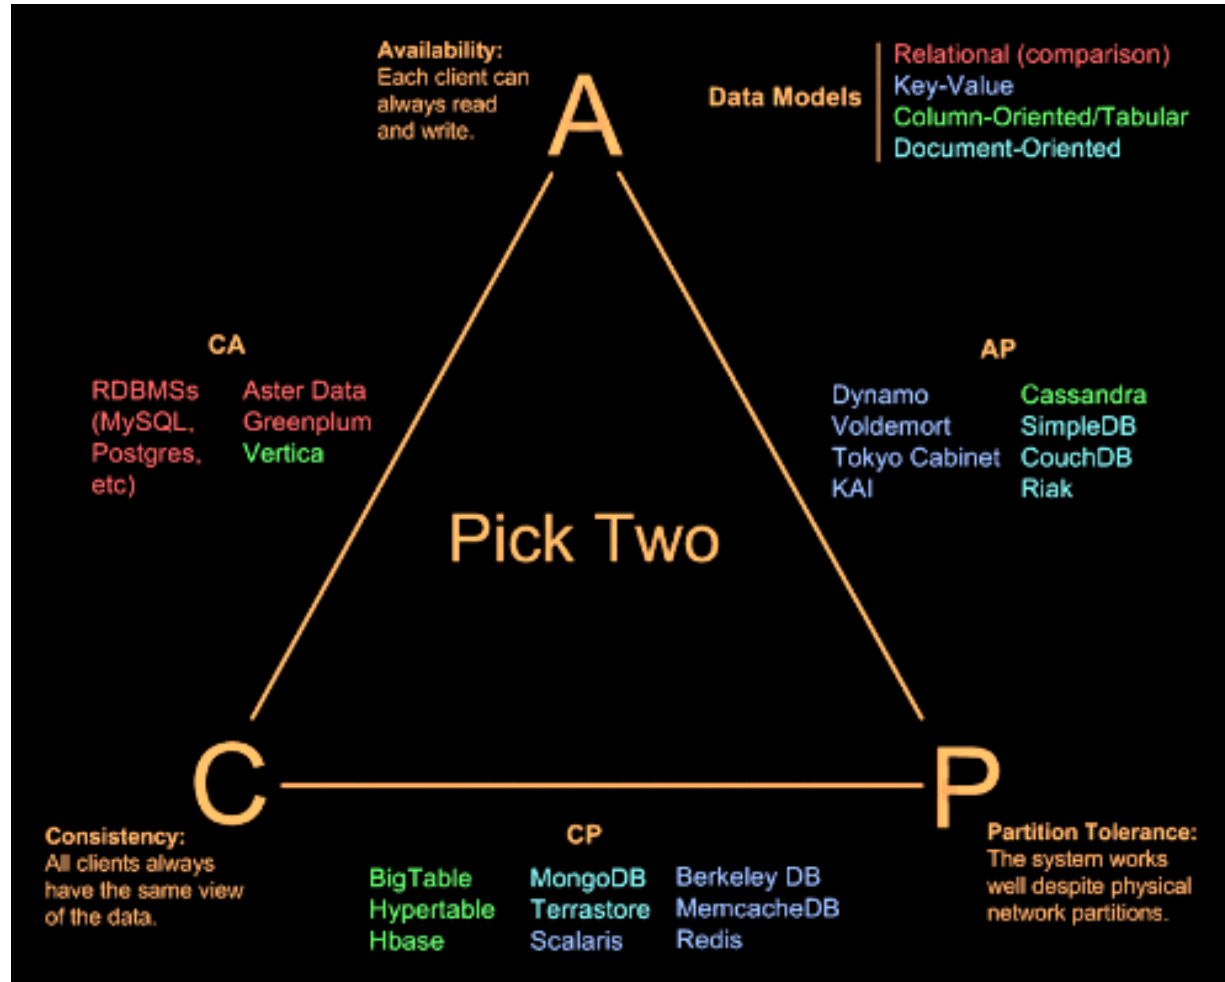
It's a trade-off!


It's a
continuous
space!

It's not at
DBMS level,

It's at
operation level!

Visual Guide to CAP Theorem





ACID vs. BASE properties

Recall the CAP theorem!

No transactional properties?!

- ACID properties may not hold -> no properties at all then???
- focuses on availability of data even in the presence of multiple failures
- spread data across many storage systems with a high degree of replication

ACID vs. BASE

- **Rationale:**

- It's ok to use stale data (Accounting systems do this all the time. It's called "closing out the books.") ; it's ok to give approximate answers
- Use resource versioning -> say what the data really is about – no more, no less.
- The value of x is 5, at time T and date D
- - So, shift the PH from 0-6 (acidic) to 8-14 (basic) – pure water's PH is 7 and neutral

- **Can some compromise be made between C and A?:**

- instead of completely giving up on C, for A
- Instead of completely giving up on A, instead of C

BASE

(Basically Available, Soft-State, Eventually Consistent)

- **Basic Availability:** fulfill request, even in partial consistency.
- **Soft State:** abandon the consistency requirements of the ACID model pretty much completely
- **Eventual Consistency:** at some point in the future, data will converge to a consistent state; delayed consistency, as opposed to immediate consistency of the ACID properties.
 - purely a [liveness](#) guarantee (reads eventually return the requested value); but
 - does not make [safety](#) guarantees, i.e.,
 - an eventually consistent system can return any value before it converges

ACID vs. BASE trade-off

- No general answer to whether your application needs an ACID versus BASE consistency model.
- Given **BASE's** loose consistency, developers need to be more **knowledgeable and rigorous about consistent data** if they choose a BASE store for their application.
- Planning around BASE limitations can sometimes be a major disadvantage when compared to the simplicity of ACID transactions.
- A fully ACID database is the perfect fit for use cases where data reliability and consistency are essential.

Why didn't we change Relational?

- Consolidated
- Integration platform
- Granting ACID properties!

The NOSQL World

The path to the Big Data Era

Good Quote!

- *Google, Amazon, Facebook, and DARPA all recognized that **when you scale systems large enough, you can never put enough iron in one place to get the job done** (and you wouldn't want to, to prevent a single point of failure).*
- *Once you accept that you **have a distributed system, you need to give up consistency or availability**, which the fundamental transactionality of traditional RDBMSs cannot abide.*
- - Cedric Beust

NoSQL History

- MultiValue databases at TRW in 1965.
- DBM is released by AT&T in 1979.
- Lotus Domino released in 1989.
- Carlo Strozzi used the term **NoSQL** in 1998 to name his lightweight, open-source relational database that did not expose the standard SQL interface.
- Graph database Neo4j is started in 2000.
- Google BigTable is started in 2004. Paper published in 2006.
- CouchDB is started in 2005.

NoSQL History

- The research paper on Amazon Dynamo is released in 2007.
- The document database MongoDB is started in 2007 as a part of an open
- source cloud computing stack and first standalone release in 2009.
- Facebooks open sources the Cassandra project in 2008.
- Project Voldemort started in 2008.
- The term **NoSQL** was reintroduced in early 2009.

The #NOSQL story

- A meetup
 - Johan Oskarsson in SF, CA, June 2009
- A hashtag: #nosql

Kinds of NoSQL

- NoSQL solutions fall into two major areas:
 - **Key/Value** or 'the big hash table' (remember caching?)
 - Amazon S3 (Dynamo)
 - Voldemort
 - Scalaris
 - MemcacheDB,
 - Azure Table Storage,
 - Redis,
 - Riak
 - **Schema-less**
 - Cassandra (column-based)
 - CouchDB (document-based)
 - Neo4J (graph-based)
 - HBase (column-based)

Different Types of NoSQL

- **Key-Value Store**

- A key that refers to a payload (actual content / data)
- MemcacheDB, Azure Table Storage, Redis

- **Column Store**

- Column data is saved together, as opposed to row data
- Super useful for data analytics
- Hadoop, Cassandra, Hypertable

Different Types of NoSQL

- **Document / XML / Object Store**

- Key (and possibly other indexes) point at a serialized object
- DB can operate against values in document
- MongoDB, CouchDB, RavenDB

- **Graph Store**

- Nodes are stored independently, and the relationship between nodes (edges) are stored with data
- Neo4j

NoSQL originators

- Google (BigTable, LevelDB)
- LinkedIn (Voldemort)
- Facebook (Cassandra)
- Twitter (Hadoop/Hbase, FlockDB, Cassandra)
- Netflix (SimpleDB, Hadoop/HBase, Cassandra)
- CERN (CouchDB)

Transactional properties?!

- **SQL databases:**


- Structured query language for
- Traditional relational databases (unique keys, single valued, no update/insertion/deletion anomalies)
- Well structured data
- ACID properties should hold

- **NoSQL (Not only SQL) databases:**


- triggered by the storage needs of [Web 2.0](#) companies such as [Facebook](#), [Google](#) and [Amazon.com](#)
- Not necessarily well structured – e.g., pictures, documents, web page description, video clips, etc.
- Lately of increasing importance due to big data
- BASE properties may be enough

Integration: from data to services

SQL
world



NO-SQL
world





POLITECNICO
MILANO 1863

SYSTEMS AND METHODS FOR BIG AND UNSTRUCTURED DATA

Towards NoSQL

Marco Brambilla

marco.brambilla@polimi.it

 @marcobrambi