



**POLITECNICO**  
MILANO 1863

SYSTEMS AND METHODS FOR BIG AND UNSTRUCTURED DATA

# SPARK

Marco Brambilla

marco.brambilla@polimi.it

 @marcobrambi

# What is Big Data used For?

Reports, e.g.,

- Track business processes, transactions

Diagnosis, e.g.,

- Why is user engagement dropping?

- Why is the system slow?

- Detect spam, worms, viruses, DDoS attacks

Decisions, e.g.,

- Personalized medical treatment

- Decide what feature to add to a product

- Decide what ads to show

Data is only as useful as the decisions it enables

# Data Processing Goals



**Low latency (interactive) queries on historical data:** enable faster decisions

E.g., identify why a site is slow and fix it



**Low latency queries on live data (streaming):** enable decisions on real-time data

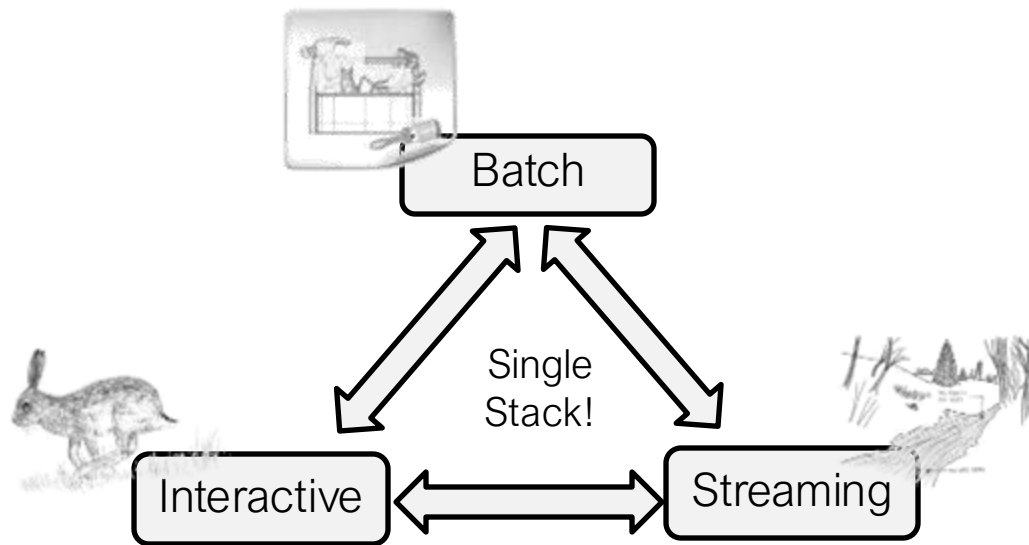
E.g., detect & block worms in real-time (a worm may infect **1mil** hosts in **1.3sec**)



**Sophisticated data processing:** enable “better” decisions

E.g., anomaly detection, trend analysis

# Goal



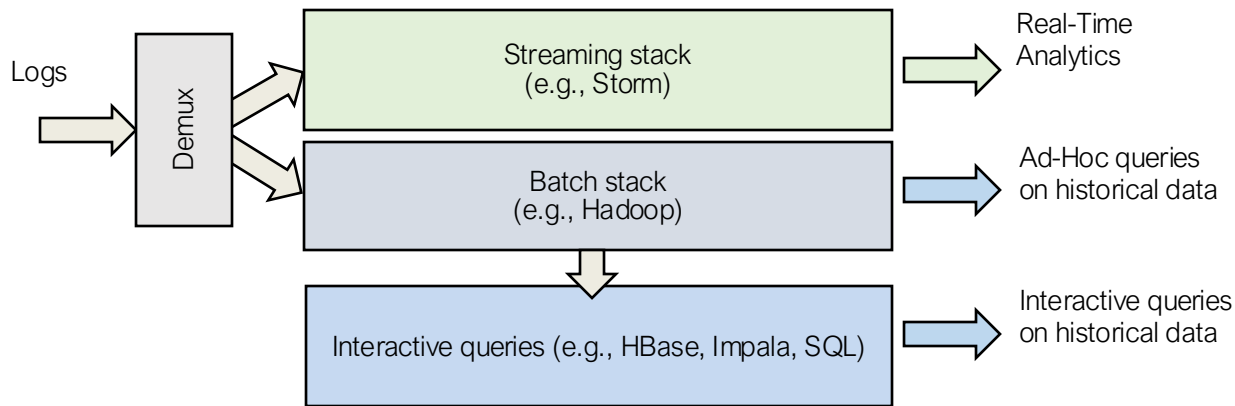
Support *batch*, *streaming*, and *interactive* computations...

... and make it easy to compose them

*Easy* to develop *sophisticated* algorithms (e.g., graph, ML algos)

# The Need for Unification (1/2)

## Today's state-of-art analytics stack



## Challenges:

- » Need to maintain three separate stacks
  - Expensive and complex
  - Hard to compute consistent metrics across stacks
- » Hard and slow to share data across stacks

# The Need for Unification (2/2)

Make real-time decisions

Detect DDoS, fraud, etc

E.g.,: what's needed to detect a DDoS attack?

1. Detect attack pattern in real time → streaming
2. Is traffic surge expected? → interactive queries
3. Making queries fast → pre-computation (batch)

And need to implement complex algos (e.g., ML)!



# Data Processing Stack



Data Processing Layer

Resource Management Layer

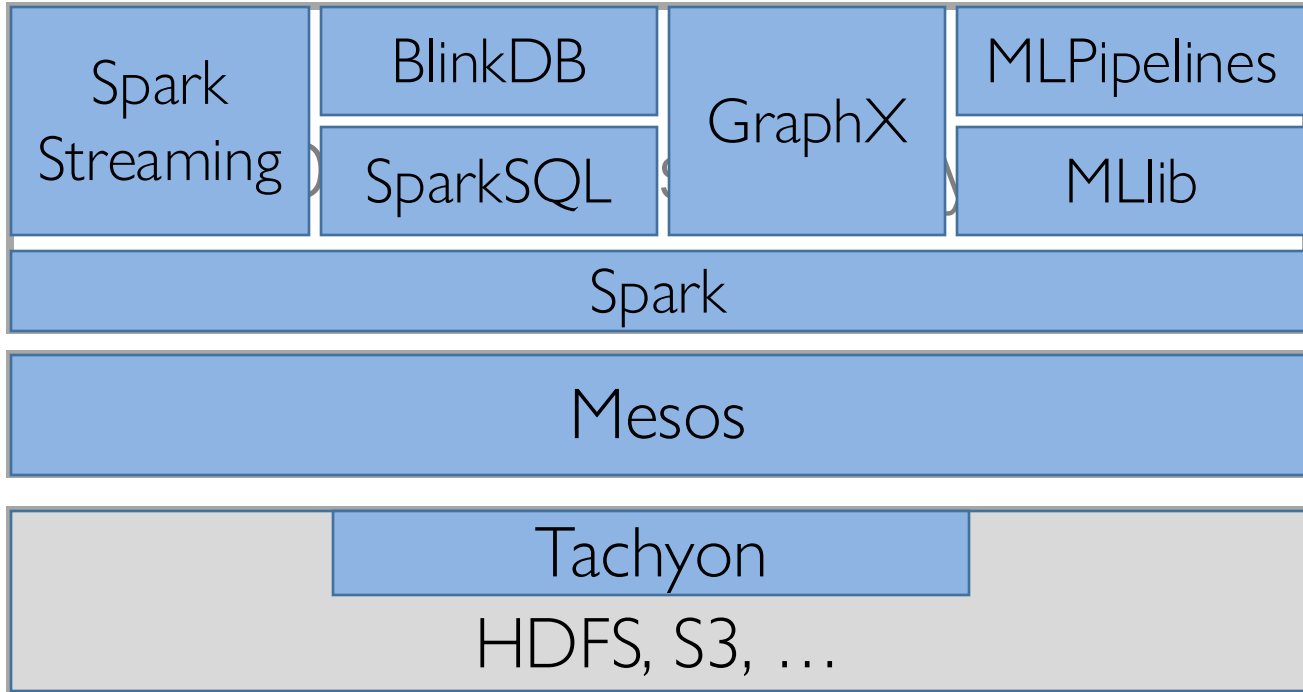
Storage Layer

# Spark

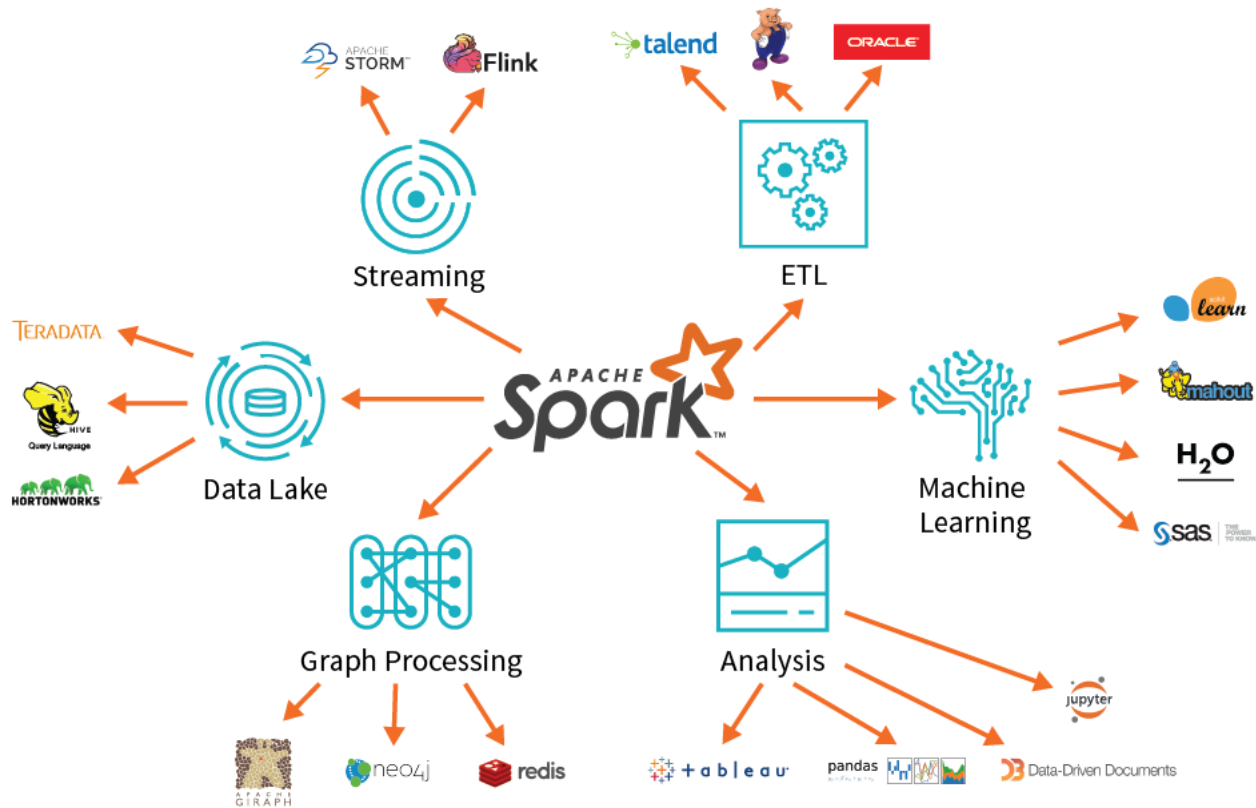
- An open source project on Apache
- First released in February 2013 and has exploded in popularity due to its ease of use and speed
- It was created at the AMPLab at UC Berkeley



# Berkeley Data Analytics Stack



# Spark



# HDFS

# HDFS Overview

Responsible for storing data on the cluster

Data files are split into blocks and distributed across the nodes in the cluster

Each block is replicated multiple times

# HDFS Basic Concepts

HDFS is a file system written in Java based on the Google's GFS

Provides redundant storage for massive amounts of data

# HDFS Basic Concepts

HDFS works best with a smaller number of large files

- Millions as opposed to billions of files

- Typically 100MB or more per file

Files in HDFS are write once

Optimized for streaming reads of large files and not random reads

# How are Files Stored

Files are split into blocks

Blocks are split across many machines at load time

Different blocks from the same file will be stored on different machines

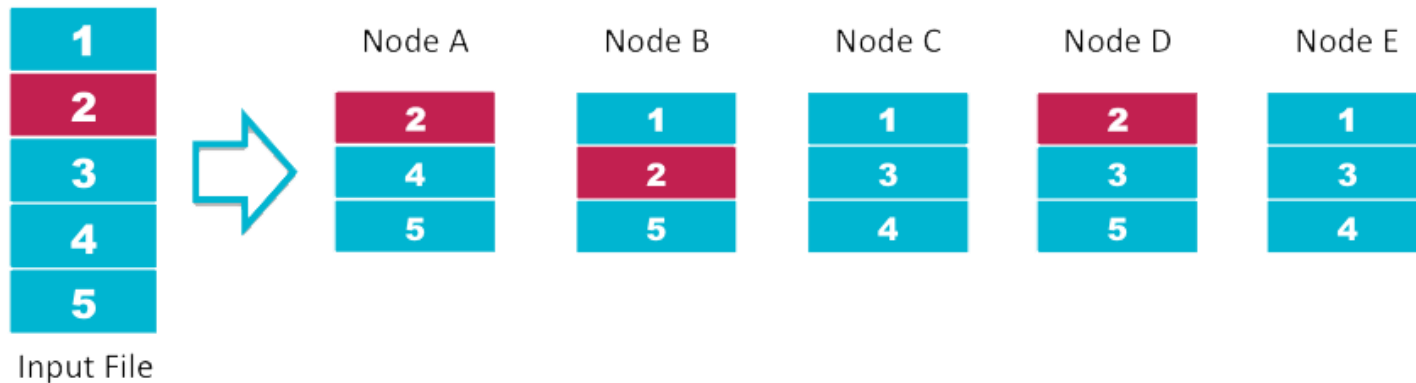
Blocks are replicated across multiple machines

The NameNode keeps track of which blocks make up a file and where they are stored

# Data Replication

Default replication is 3-fold

HDFS Data Distribution





# Goals of HDFS



Very Large Distributed File System

10K nodes, 100 million files, 10PB

Assumes Commodity Hardware

Files are replicated to handle hardware failure

Detect failures and recover from them

Optimized for Batch Processing

Data locations exposed so that computations can move to where data resides

Provides very high aggregate bandwidth

# Distributed File System

**Single Namespace** for entire cluster

**Data Coherency**

- Write-once-read-many** access model

- Client can only append to existing files

**Files are broken up into blocks**

- Typically 64MB block size

- Each block replicated on multiple DataNodes

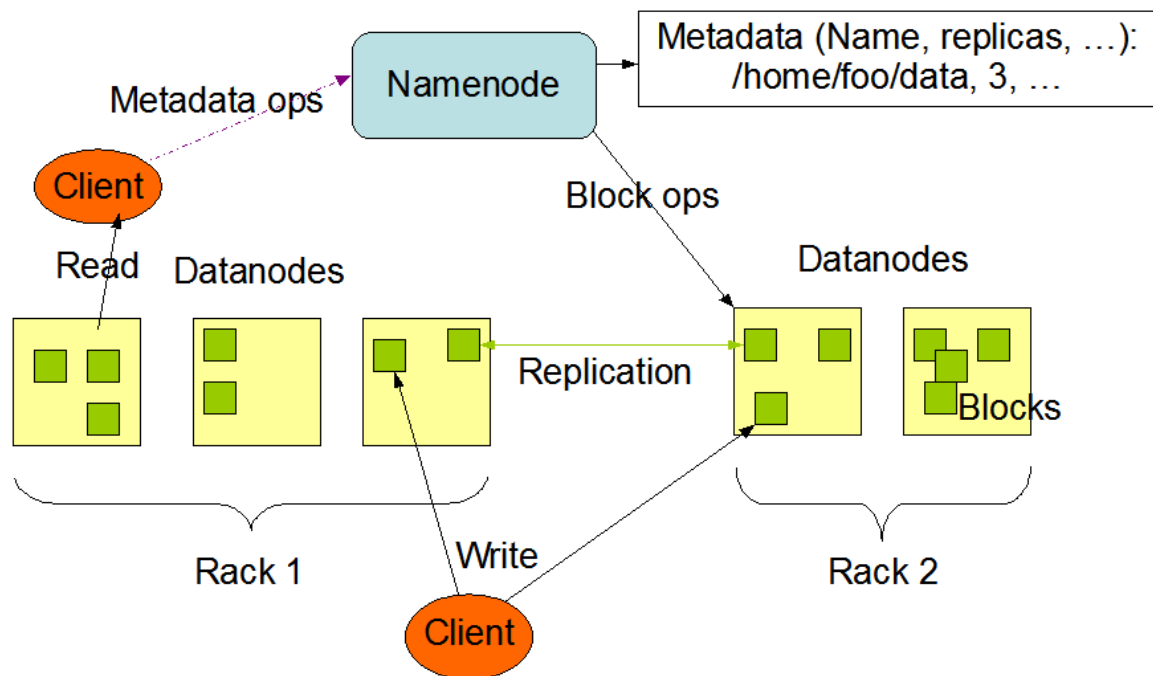
**Intelligent Client**

- Client can find location of blocks

- Client accesses data directly from DataNode

# HDFS Architecture

HDFS Architecture



# Functions of a NameNode

Manages File System Namespace

- Maps a file name to a set of blocks

- Maps a block to the DataNodes where it resides

Cluster Configuration Management

Replication Engine for Blocks

# NameNode Metadata

## Metadata in Memory

- The entire metadata is in main memory

- No demand paging of metadata

## Types of metadata

- List of files

- List of Blocks for each file

- List of DataNodes for each block

- File attributes, e.g. creation time, replication factor

## A Transaction Log

- Records file creations, file deletions etc

# DataNode

## A Block Server

- Stores data in the local file system (e.g. ext3)

- Stores metadata of a block (e.g. CRC)

- Serves data and metadata to Clients

## Block Report

- Periodically sends a report of all existing blocks to the NameNode

## Facilitates Pipelining of Data

- Forwards data to other specified DataNodes

# Block Placement

## Strategy

- One replica on local node

- Second replica on a remote rack

- Third replica on same remote rack

- Additional replicas are randomly placed

Clients read from nearest replicas

# Heartbeats

DataNodes send heartbeat to the NameNode

Once every 3 seconds

NameNode uses heartbeats to detect

DataNode failure



# Replication Engine

NameNode detects DataNode failures

- Chooses new DataNodes for new replicas

- Balances disk usage

- Balances communication traffic to DataNodes

# Data Correctness

Use Checksums to validate data

- Use CRC32

File Creation

- Client computes checksum per 512 bytes

- DataNode stores the checksum

File access

- Client retrieves the data and checksum from DataNode

- If Validation fails, Client tries other replicas

# NameNode Failure

A single point of failure in HDFS 1

Transaction Log stored in multiple directories

- A directory on the local file system

- A directory on a remote file system (NFS/CIFS)

# Data Pipelining

Client retrieves a list of DataNodes on which to place replicas of a block

Client writes block to the first DataNode

The first DataNode forwards the data to the next node in the Pipeline

When all replicas are written, the Client moves on to write the next block in file

# Rebalancer

Goal: % disk full on DataNodes should be similar

- Usually run when new DataNodes are added
- Cluster is online when Rebalancer is active
- Rebalancer is throttled to avoid network congestion

# Secondary NameNode

Copies FsImage and Transaction Log from Namenode to a temporary directory

Merges FSImage and Transaction Log into a new FSImage in temporary directory

Uploads new FSImage to the NameNode

Transaction Log on NameNode is purged

# User Interface

## Commands for HDFS User:

```
hadoop dfs -mkdir /foodir  
hadoop dfs -cat /foodir/myfile.txt  
hadoop dfs -rm /foodir/myfile.txt
```

## Commands for HDFS Administrator

```
hadoop dfsadmin -report  
hadoop dfsadmin -decommission datanodename
```

## Web Interface

```
http://host:port/dfshealth.jsp
```

# Data Retrieval

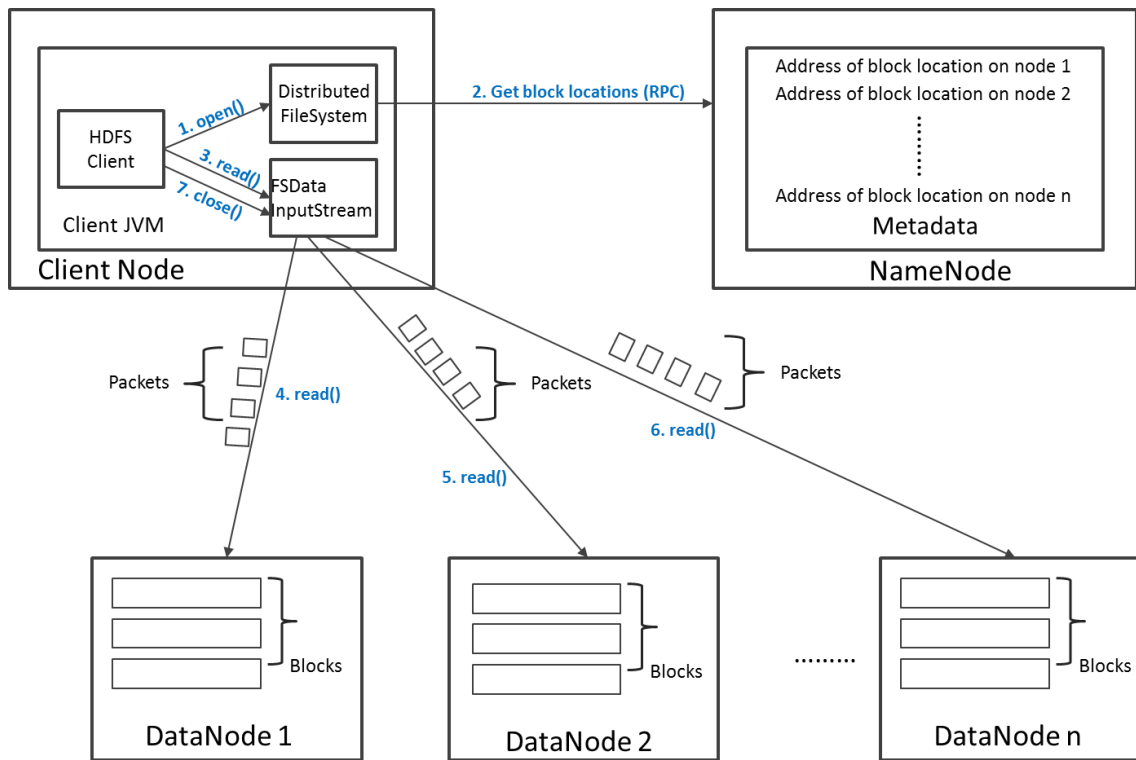
When a client wants to retrieve data

Communicates with the NameNode to determine which blocks make up a file and on which data nodes those blocks are stored

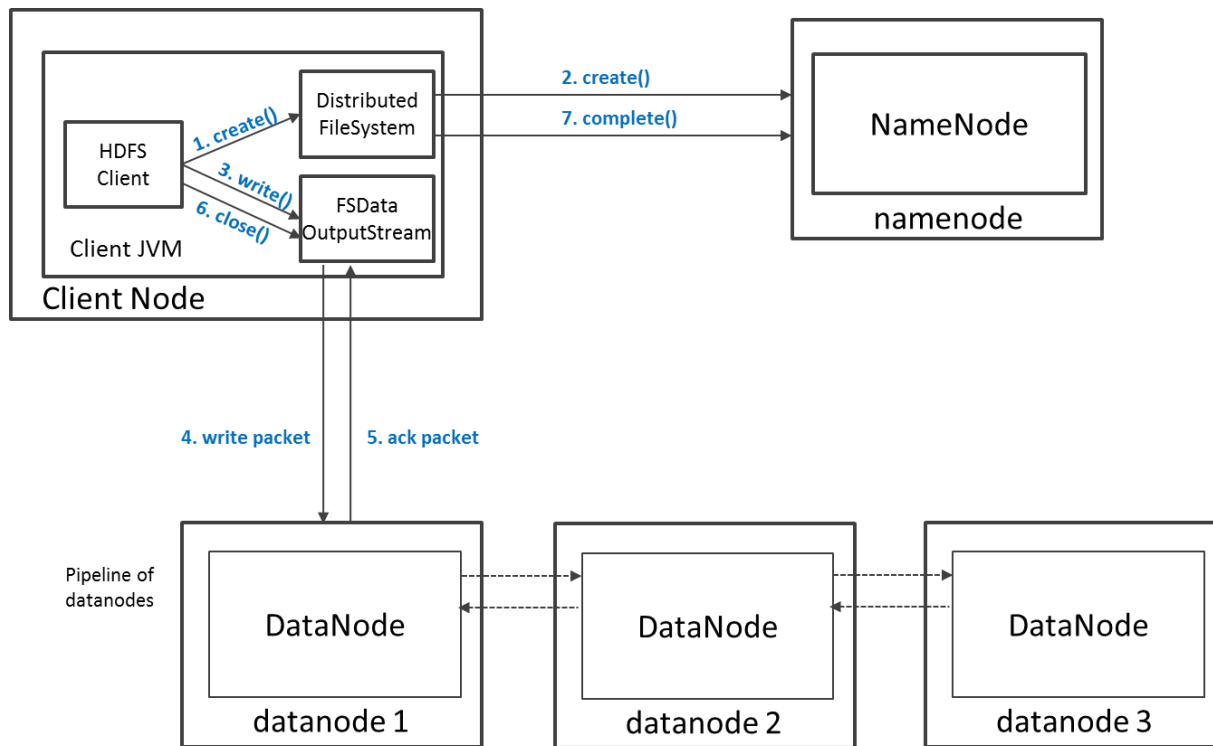
Then communicated directly with the data nodes to read the data



# Read Operation in HDFS



# Write Operation in HDFS



# HDFS Security

## Authentication to Hadoop

Simple – insecure way of using OS username to determine hadoop identity

Kerberos – authentication using kerberos ticket

Set by `hadoop.security.authentication=simple|kerberos`

## File and Directory permissions are same like in POSIX

read (r), write (w), and execute (x) permissions

also has an owner, group and mode

enabled by default (`dfs.permissions.enabled=true`)

## ACLs are used for implementation permissions that differ from natural hierarchy of users and groups

enabled by `dfs.namenode.acls.enabled=true`

# HDFS Configuration

## HDFS Defaults

Block Size – 64 MB

Replication Factor – 3

Web UI Port – 50070

HDFS conf file – `/etc/hadoop/conf/hdfs-site.xml`

# Interfaces to HDFS

Java API (`DistributedFileSystem`)

C wrapper (`libhdfs`)

HTTP protocol

WebDAV protocol

Shell Commands

However the command line is one of the simplest and most familiar

# HDFS – Shell Commands

There are two types of shell commands

## User Commands

`hdfs dfs` – runs filesystem commands on the HDFS

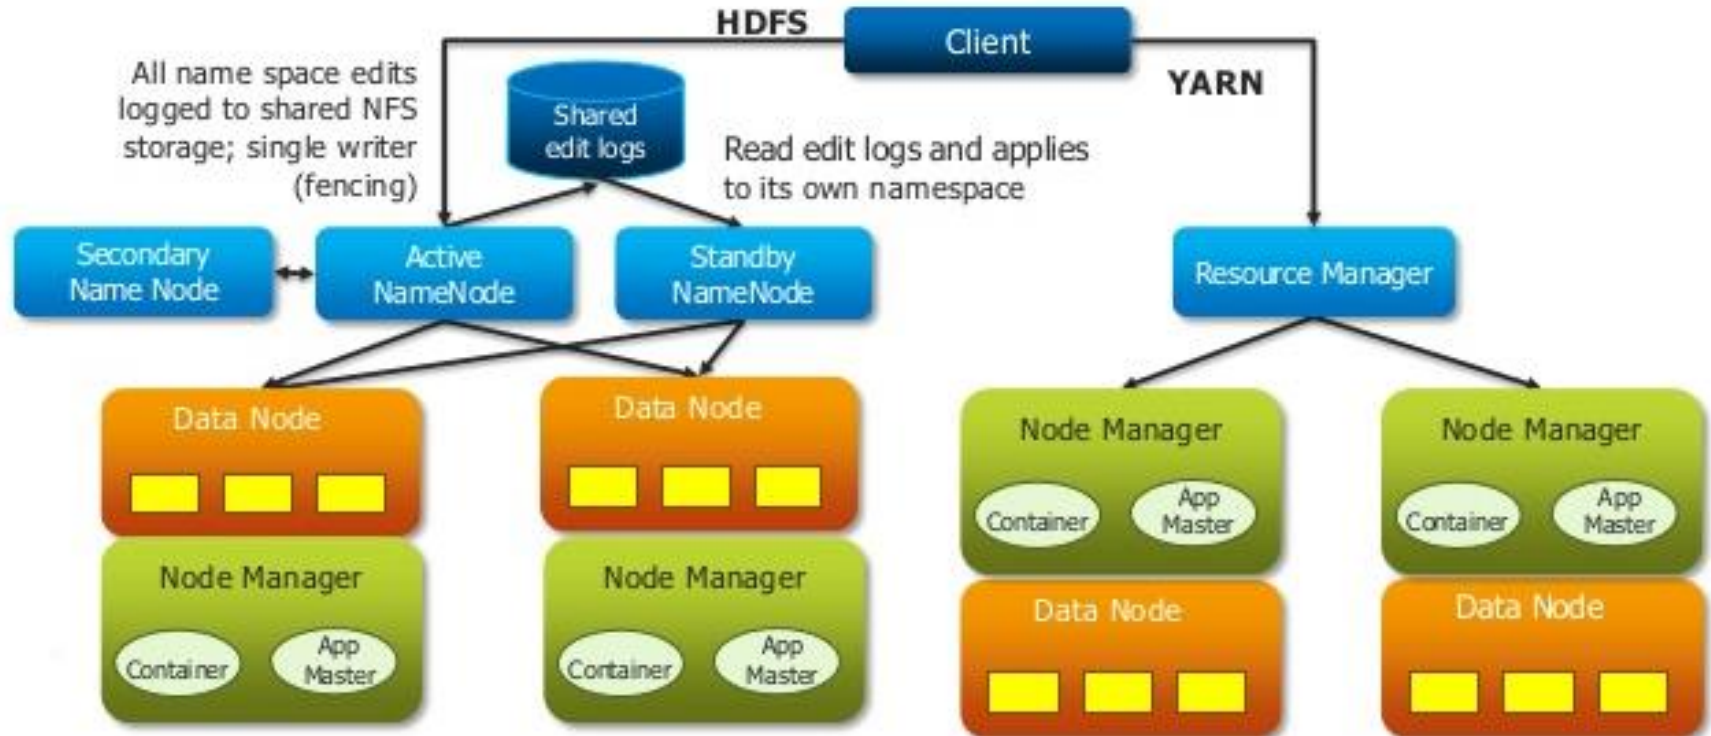
`hdfs fsck` – runs a HDFS filesystem checking command

## Administration Commands

`hdfs dfsadmin` – runs HDFS administration commands

# Hadoop

# Hadoop 2.0 Architecture – YARN





# Spark

# History

[Apache Spark](#) started as a research project at the University of California AMPLab, in 2009 by [Matei Zaharia](#).

In 2013

- donated to the Apache Software Foundation
- open sourced, adopted the Apache 2.0 license

In February 2014, Spark became a Top-Level [Apache Project](#).

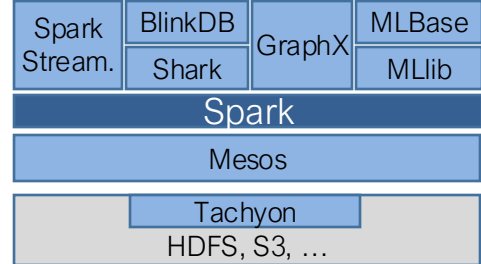
In November 2014, Spark founder Matei\_Zaharia's company [Databricks](#) set a new world record in large scale sorting using Spark.

Latest stable release: [CLICK-HERE](#)

600,000+ lines of code (75% Scala)

Built by 1,000+ developers from more than 250+ organizations

# Apache Spark



## Distributed Execution Engine

Fault-tolerant, efficient in-memory storage  
(RDDs)

Powerful programming model and APIs (Scala,  
Python, Java)

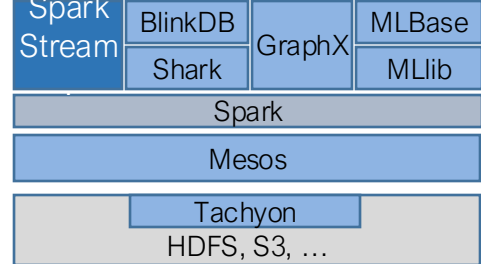
**Fast:** up to 100x faster than Hadoop

**Easy to use:** 5-10x less code than Hadoop

**General:** support interactive & iterative  
apps

Two major releases since last AMPCamp

# Spark Streaming



Large scale streaming computation  
Implement streaming as a sequence  
of <1s jobs

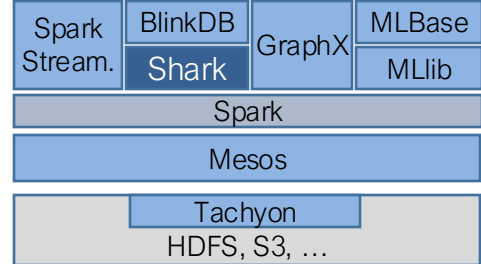
- Fault tolerant

- Handle stragglers

- Ensure exactly one semantics

Integrated with Spark: unifies **batch**,  
**interactive**, and **batch** computations

# SparkSQL



Hive over Spark: full support for HQL and UDFs

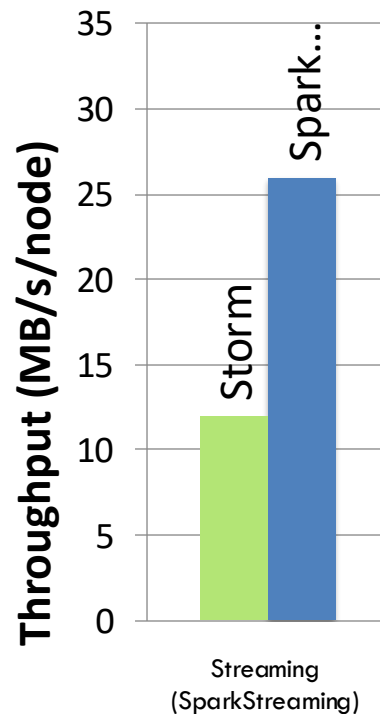
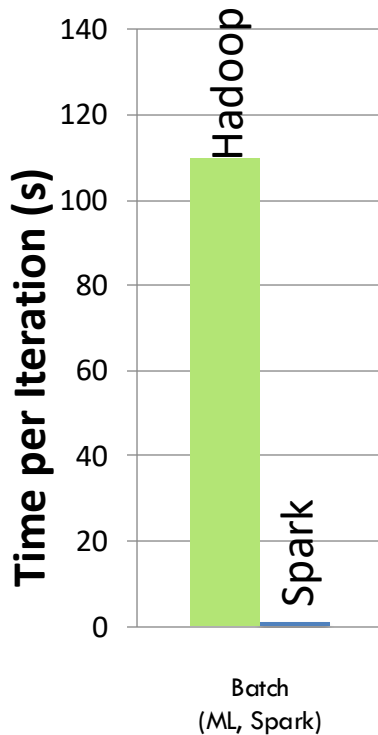
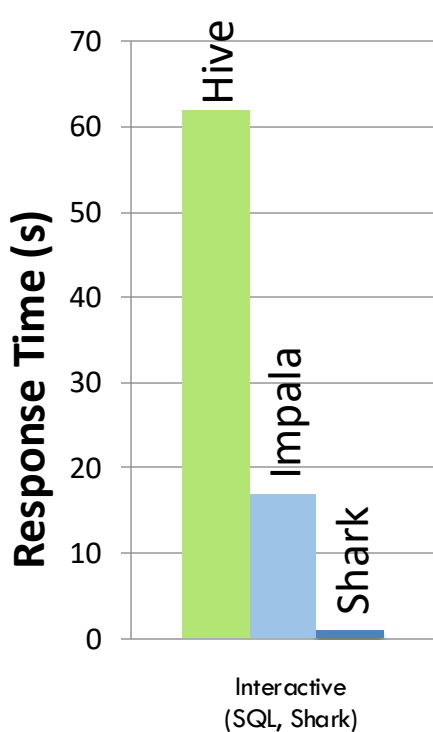
Up to 100x when input is in memory

Up to 5-10x when input is on disk

Running on hundreds of nodes at Yahoo!

Two major releases along Spark

# Performance and Generality (Unified Computation Models)



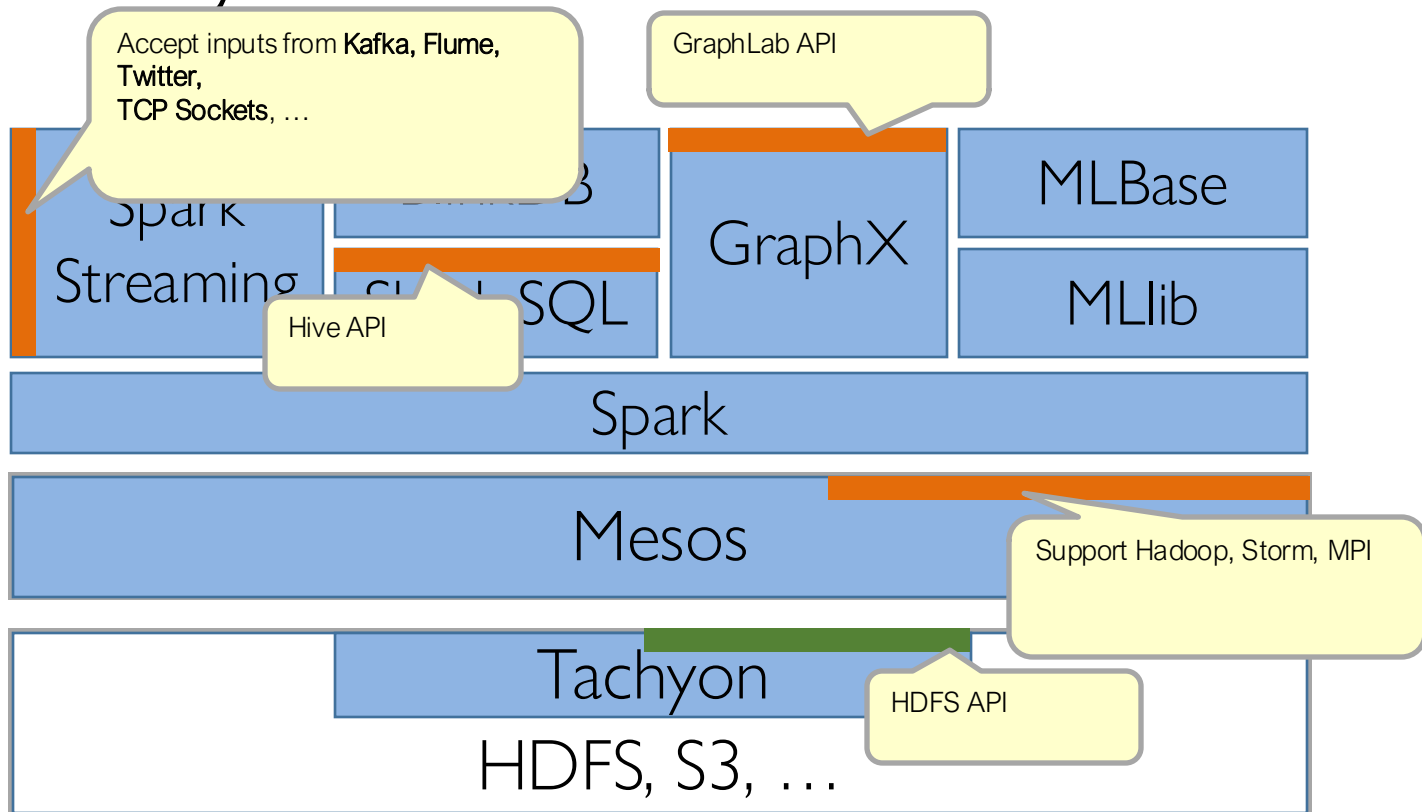
# Unified Programming Models

Unified system for  
SQL, graph  
processing,  
machine learning

All share the same  
set of workers and  
caches

```
def logRegress(points: RDD[Point]): Vector {  
  var w = Vector(D, _ => 2 * rand.nextDouble - 1)  
  for (i <- 1 to ITERATIONS) {  
    val gradient = points.map { p =>  
      val denom = 1 + exp(-p.y * (w dot p.x))  
      (1 / denom - 1) * p.y * p.x  
    }.reduce(_ + _)  
    w -= gradient  
  }  
  w  
}  
  
val users = sql2rdd("SELECT * FROM user u  
  JOIN comment c ON c.uid=u.uid")  
  
val features = users.mapRows { row =>  
  new Vector(extractFeature1(row.getInt("age")),  
    extractFeature2(row.getStr("country")),  
    ...) }  
  
val trainedVector = logRegress(features.cache())
```

# Compatibility to Existing Ecosystem

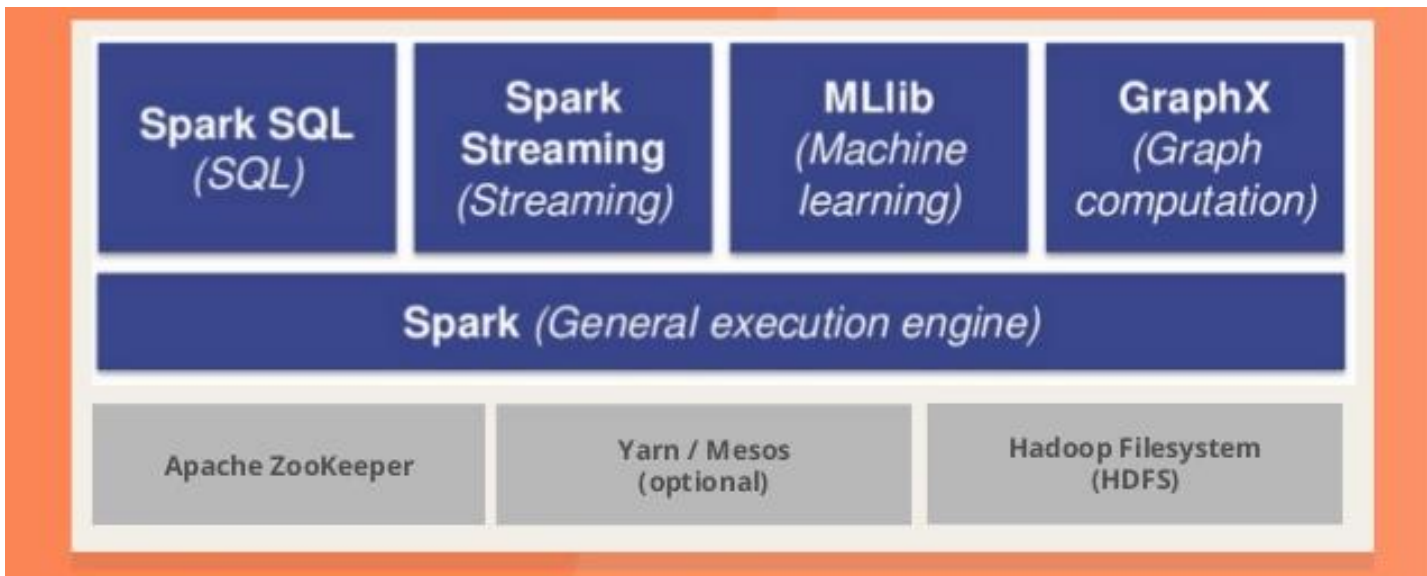




# Spark

- You can think of Spark as **scalable computation platform**
- Spark can use data stored in a variety of formats
  - Cassandra
  - AWS S3
  - HDFS
  - And more

# SPARK in context

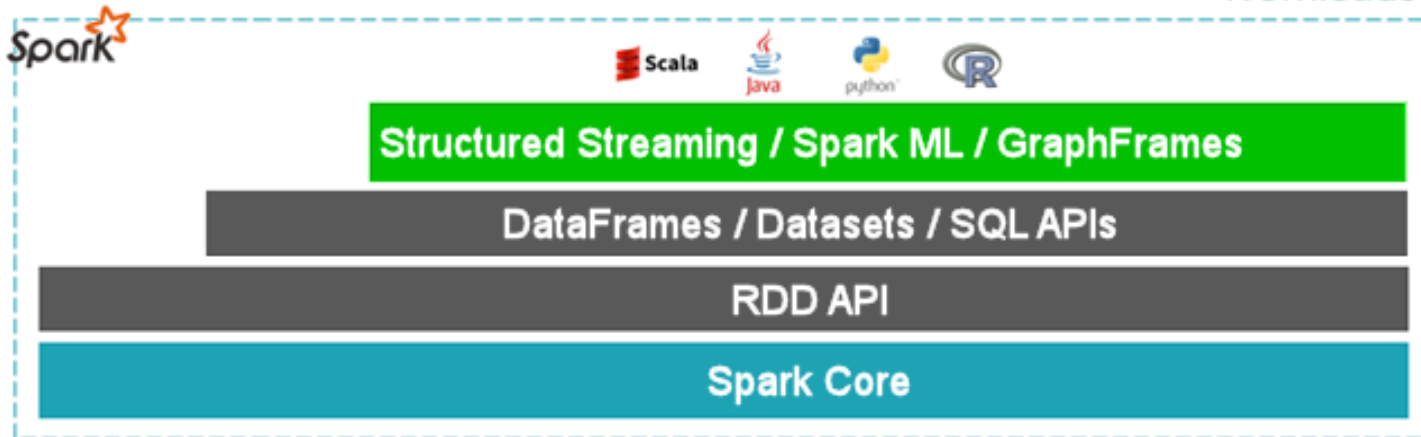


# Spark Components

## Environments



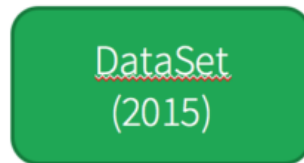
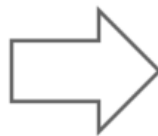
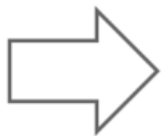
## Workloads



## Data Sources



# History of Spark APIs



Distribute collection  
of JVM objects

Functional Operators (map,  
filter, etc.)

Distribute collection  
of Row objects

Expression-based operations  
and UDFs

Logical plans and optimizer

Fast/efficient internal  
representations

Internally rows, externally  
JVM objects

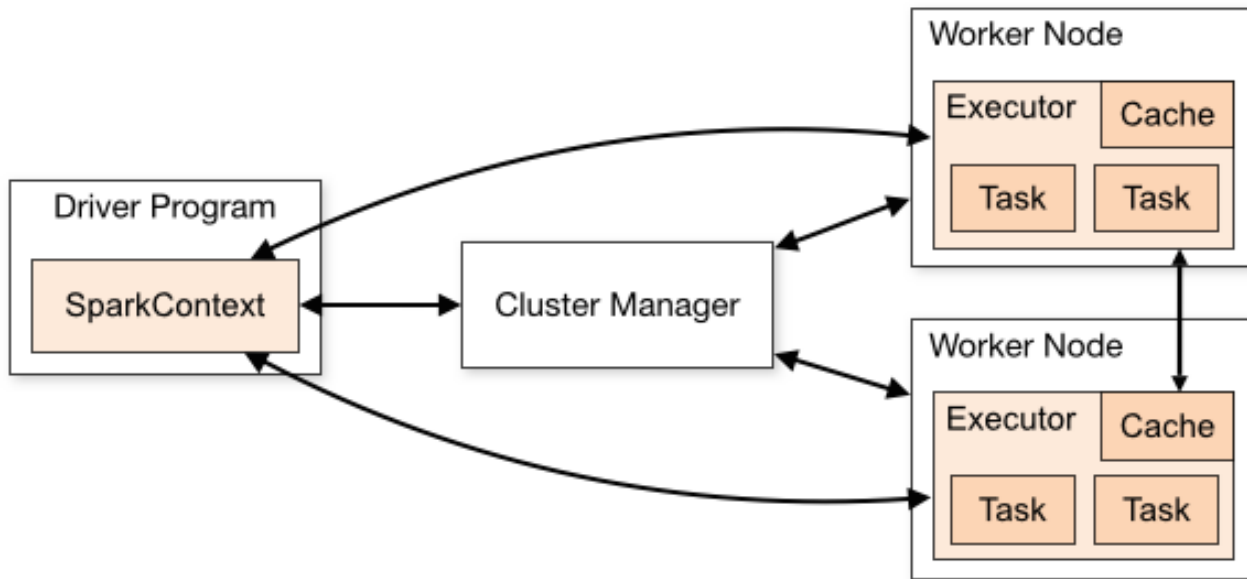
Almost the “Best of both  
worlds”: **type safe + fast**

But slower than DF  
Not as good for interactive  
analysis, especially Python

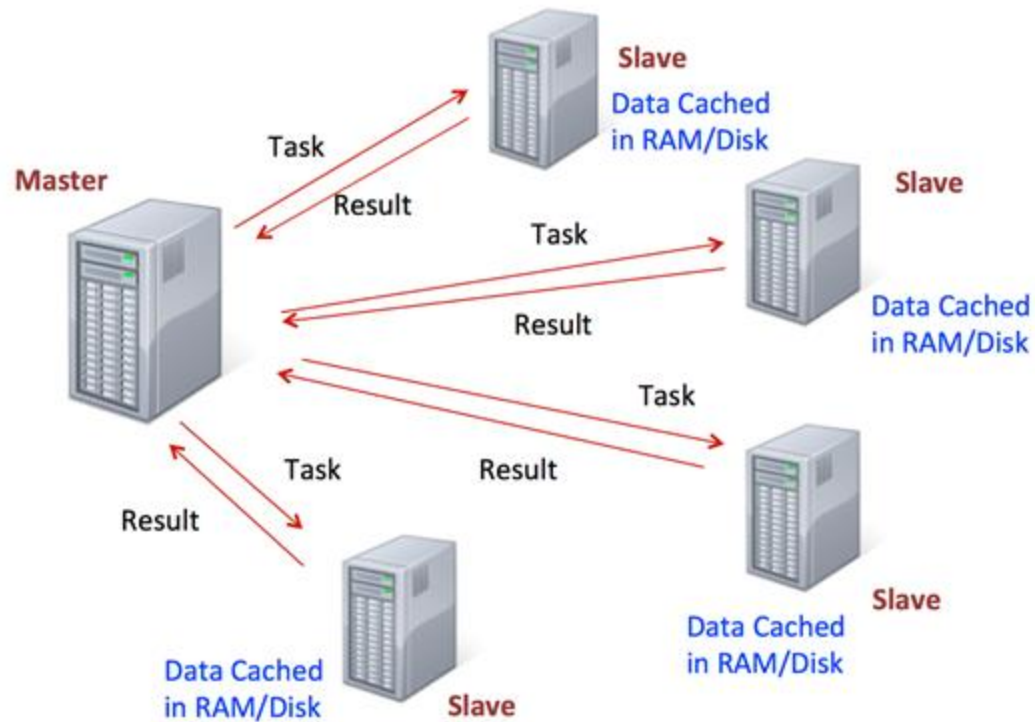
# Spark RDDs

- At the core of Spark is the idea of a Resilient Distributed Dataset (RDD)
- Resilient Distributed Dataset (RDD) has 4 main features:
  - Distributed Collection of Data
  - Fault-tolerant
  - Parallel operation - partitioned
  - Ability to use many data sources

# Spark RDDs



# Spark RDDs



# Spark RDDs

RDDs are:

- immutable,
- lazily evaluated,
- cacheable.



# Spark Operations

- There are two types of Spark operations:
  - Transformations
  - Actions
- **Transformations** are a recipe to follow.
- **Actions** perform what the recipe says to do and returns something back.



 Operations =

+



ACTIONS

# Spark Operations

- This carries over to the syntax when coding.
- you write a method call, but won't see anything as a result until you call the action.
- Why?

**With a large dataset, you don't want to calculate all the transformations until you are sure you want to perform them!**

# RDD operations

*transformations* to build RDDs through deterministic operations on other RDDs

transformations include *map*, *filter*, *join*

lazy operation

*actions* to return value or export data

actions include *count*, *collect*, *save*

triggers execution

# Operations

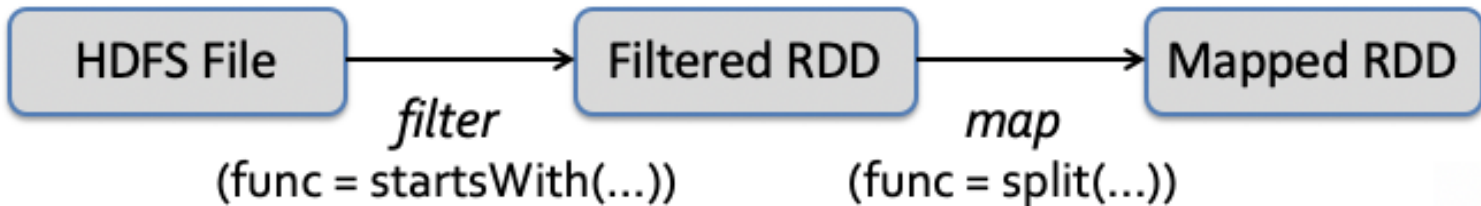


Transformations ( <i>lazy</i> )	Actions
<code>select</code>	<code>show</code>
<code>distinct</code>	<code>count</code>
<code>groupBy</code>	<code>collect</code>
<code>sum</code>	<code>save</code>
<code>orderBy</code>	
<code>filter</code>	
<code>limit</code>	

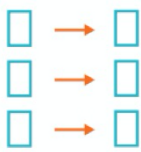

# Transformations

RDD are immutable, transformations create new RDD

```
msgs = textFile.filter(lambda s: s.startsWith("ERROR")) .map(lambda s: s.split("\t")[2])
```



# Narrow vs. Wide

Narrow Transformations	Wide Transformations
<p>The data required to compute the records in a single partition reside in at most one partition of the parent RDD.</p>	<p>The data required to compute the records in a single partition may reside in many partitions of the parent RDD.</p>
<p>Narrow Transformations 1 to 1</p>  <p>The diagram shows three blue rectangular boxes on the left representing parent partitions. Three orange arrows point from each parent box to a corresponding child box on the right, illustrating a one-to-one mapping where each child partition only depends on its parent partition.</p>	<p>Wide Transformations (shuffles) 1 to N</p>  <p>The diagram shows three blue rectangular boxes on the left representing parent partitions. Three orange arrows originate from these parent boxes and converge on a single child box on the right. This represents a shuffle operation where data from multiple parent partitions is required to compute a single child partition.</p>
<ul style="list-style-type: none"><li>* <code>filter(..)</code></li></ul>	<ul style="list-style-type: none"><li>* <code>distinct()</code></li></ul>
<ul style="list-style-type: none"><li>* <code>drop(..)</code></li></ul>	<ul style="list-style-type: none"><li>* <code>groupBy(..).sum()</code></li></ul>
<ul style="list-style-type: none"><li>* <code>coalesce()</code></li></ul>	<ul style="list-style-type: none"><li>* <code>repartition(n)</code></li></ul>

# Spark RDDs and DataFrames

- With Spark you can find: RDD syntax versus DataFrame syntax discussions. Why?
- With Spark 2.0, Spark is moving towards a DataFrame based syntax
- But the way files are being distributed can still be thought of as RDDs, it is just the typed out syntax that is changing

# DataFrame

## Definition

- Immutable Data with named columns (built on RDDs)

## Characteristics

- User-friendly API
- Uniform APIs across languages (Scala, Java, Python, R, and SQL)
- Improved performance via optimizations (Tungsten and Catalyst)



# Code example

an Apache Spark code snippet using SQL and DataFrames to query and join different data sources

```
# Read JSON file and register temp view

jsonDf = context.jsonFile("s3n://...").createOrReplaceTempView("json")

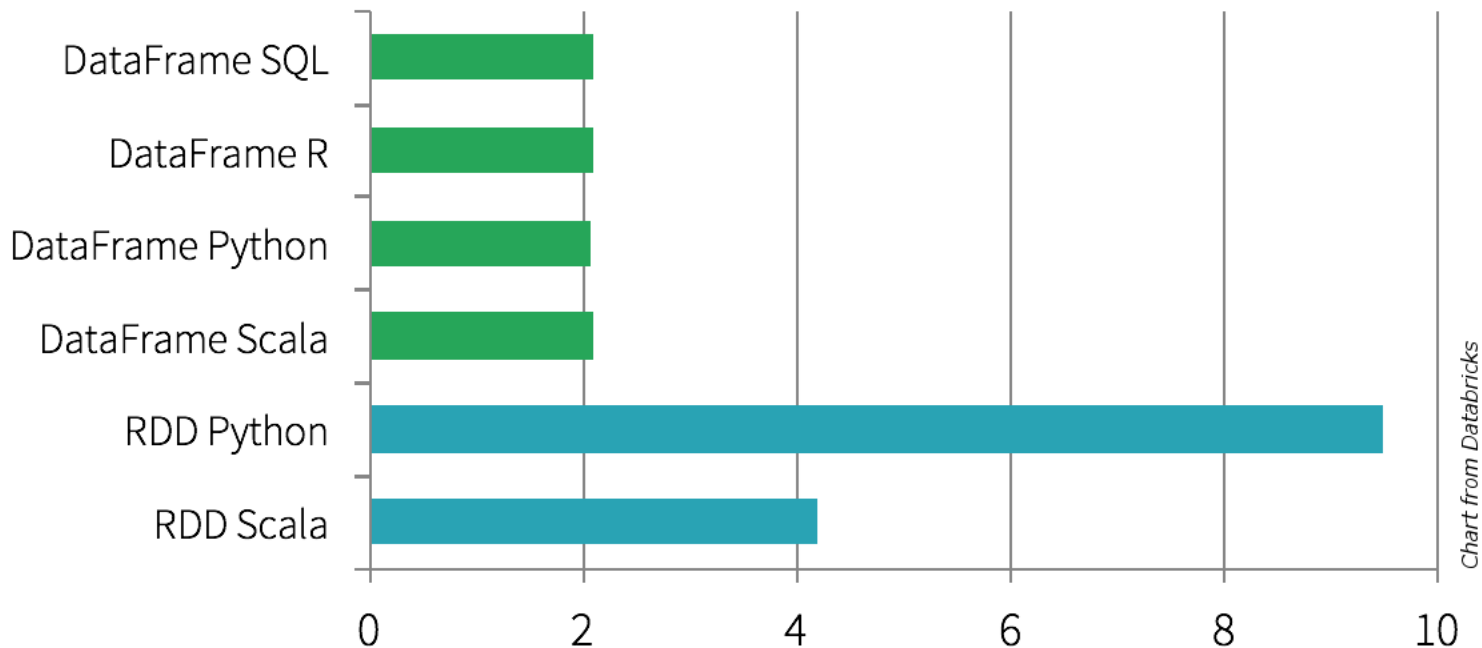
# Execute SQL query or ...

results = context.sql("""SELECT * FROM people JOIN json ... WHERE ...""")
```

```
# ... or Use DataFrame APIs

results = peopleDf.join(jsonDf, peopleDf.id == jsonDf.id).filter(...)
```

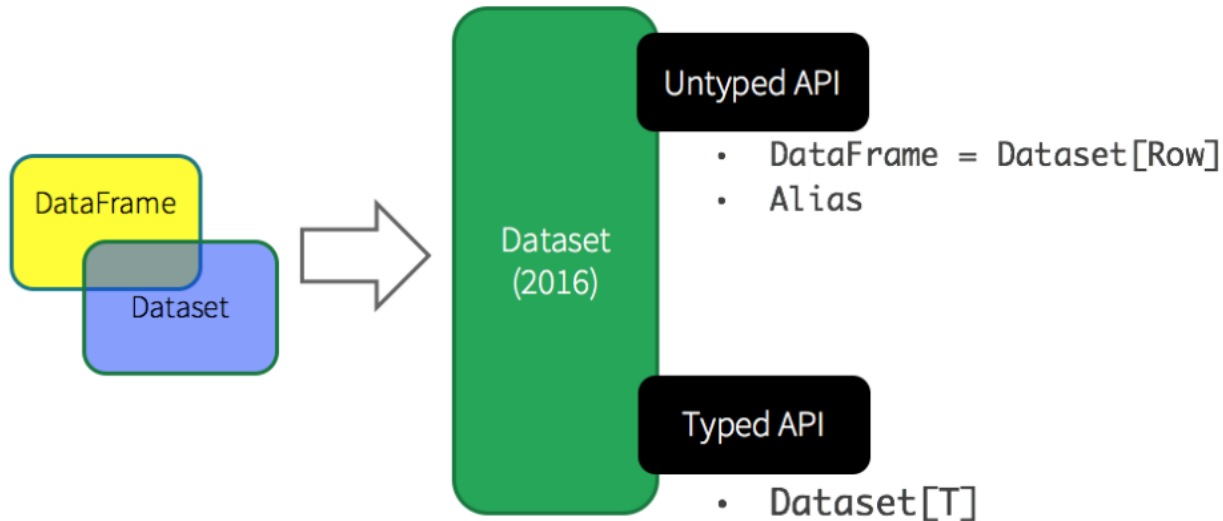
# Performance



Time to aggregate 10 million integer pairs (in seconds)

# Dataset

## Unified Apache Spark 2.0 API



# Software Components

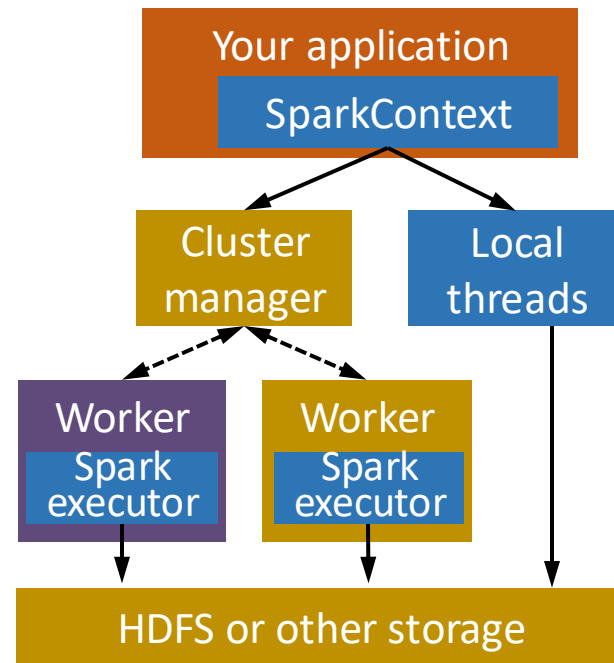
Spark runs as a library in your program (1 instance per app)

Runs tasks locally or on cluster

Mesos, YARN or standalone mode

Accesses storage systems via Hadoop InputFormat API

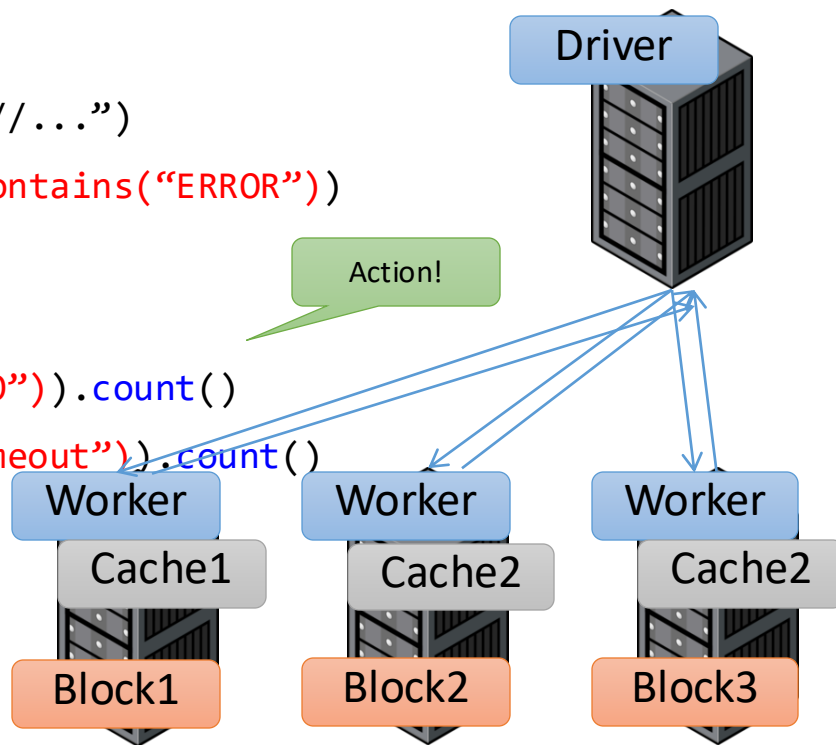
Can use HBase, HDFS, S3, ...



# Job example

```
val log = sc.textFile("hdfs://...")  
val errors = file.filter(_.contains("ERROR"))  
errors.cache()
```

```
errors.filter(_.contains("I/O")).count()  
errors.filter(_.contains("timeout")).count()
```



# Job Example

Load error messages from a log into memory, then interactively search for various patterns

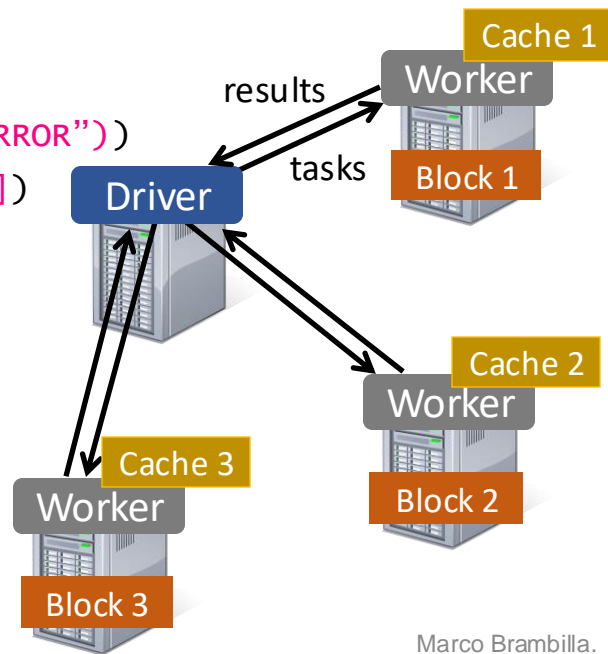
```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

...

## Full-text search of Wikipedia

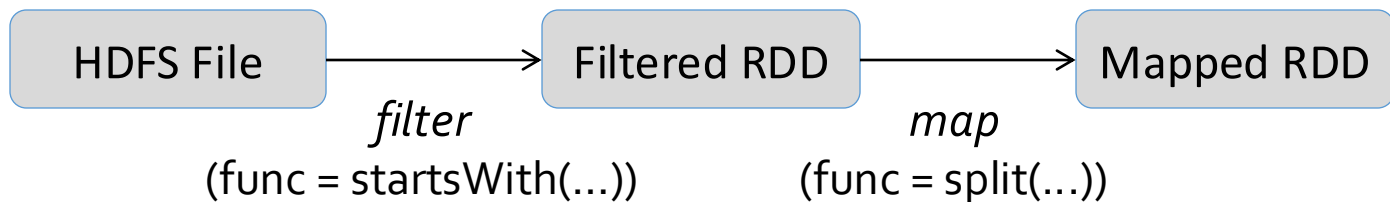
- 60GB on 20 EC2 machines
- 0.5 sec vs. 20s for on-disk



# Fault Recovery

RDDs track *lineage* information that can be used to efficiently recompute lost data

```
msgs = textFile.filter(lambda s: s.startsWith("ERROR"))  
                .map(lambda s: s.split("\t")[2])
```



# RDD partition-level view

Dataset-level view:

log:

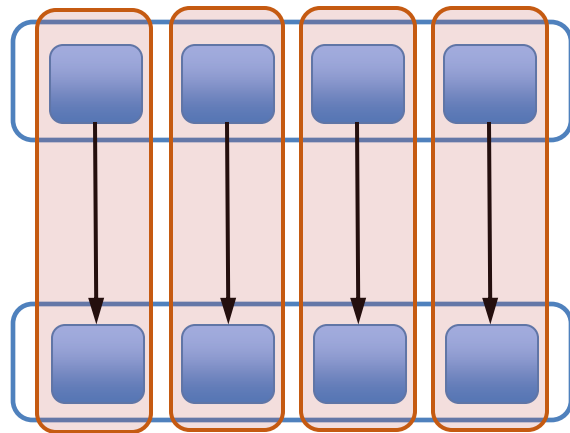
HadoopRDD  
path = hdfs://...



errors:

FilteredRDD  
func = \_.contains(...)  
shouldCache = true

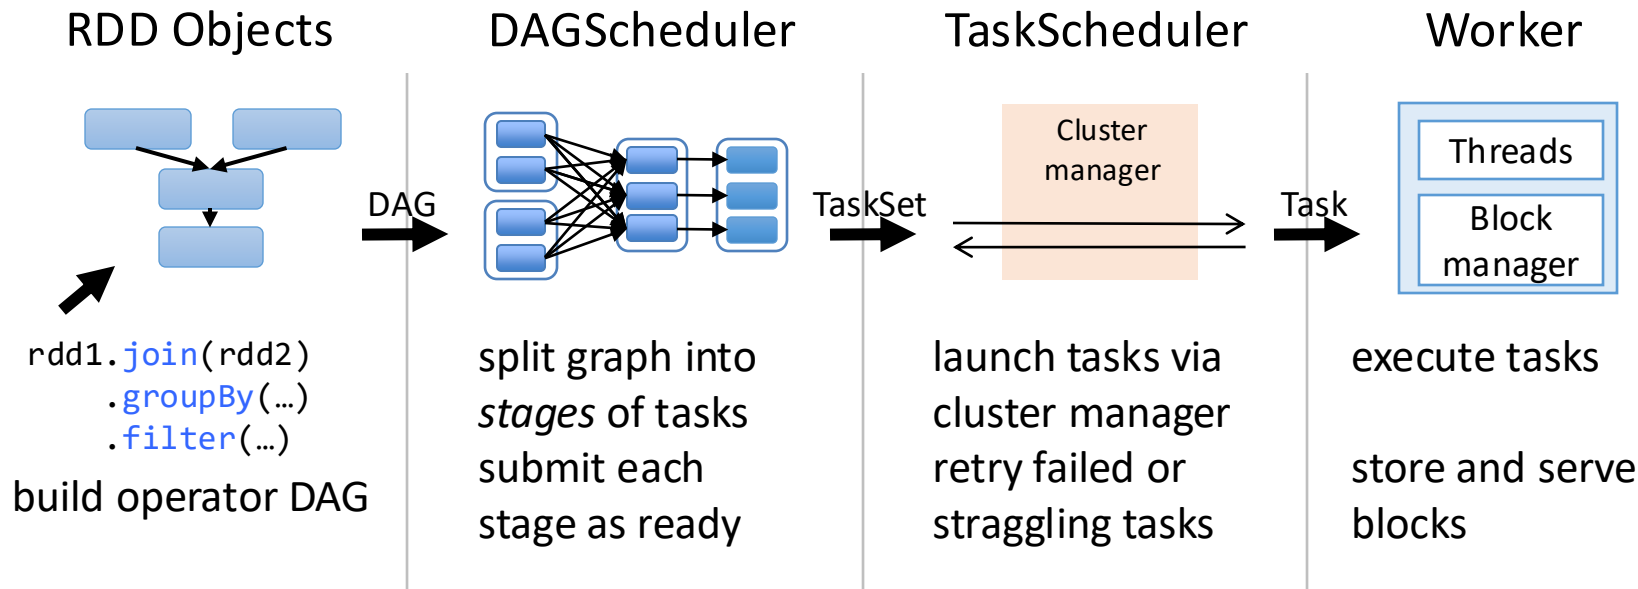
Partition-level view:



Task 1 Task 2 ...



# Job scheduling



# Available APIs

You can write in Java, Scala or Python  
interactive interpreter: Scala & Python only  
standalone applications: any  
performance: Java & Scala are faster thanks  
to static typing

# Importing Data

Importers for most  
data formats and from  
most database technology

## Spark Data Sources

- Cassandra
- Couchbase
- ElasticSearch
- Importing HIVE Tables
- MongoDB
- Neo4j
- Oracle
- Reading Avro Files
- Reading CSV Files
- Reading JSON Files
- Reading LZO Compressed Files
- Reading Parquet Files
- Redis
- Riak Time Series
- Connecting to SQL Databases using JDBC
- Zip Files
- Amazon Redshift
- Amazon S3 with Apache Spark
- Azure storage services
- Azure Cosmos DB
- SQL Data Warehouse

# SparkContext

Main entry point to Spark functionality

Available in shell as variable **SC**

In standalone programs, you'd make your own (see later for details)

# Creating RDDs

# Turn a Python collection into an RDD

```
>sc.parallelize([1, 2, 3])
```

# Load text file from local FS, HDFS, or S3

```
>sc.textFile("file.txt")
```

```
>sc.textFile("directory/*.txt")
```

```
>sc.textFile("hdfs://namenode:9000/path/file")
```

# Use existing Hadoop InputFormat (Java/Scala only)

```
>sc.hadoopFile(keyClass, valClass, inputFmt, conf)
```

# Basic Transformations

```
> nums = sc.parallelize([1, 2, 3])
```

```
# Pass each element through a function
```

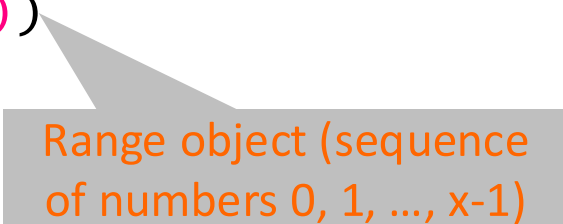
```
> squares = nums.map(lambda x: x*x)    // {1, 4, 9}
```

```
# Keep elements passing a predicate
```

```
> even = squares.filter(lambda x: x % 2 == 0) // {4}
```

```
# Map each element to zero or more others
```

```
> nums.flatMap(lambda x: => range(x))  
  ># => {0, 0, 1, 0, 1, 2}
```



Range object (sequence  
of numbers 0, 1, ..., x-1)

# Basic Actions

```
> nums = sc.parallelize([1, 2, 3])  
  
# Retrieve RDD contents as a local collection  
> nums.collect() # => [1, 2, 3]  
  
# Return first K elements  
> nums.take(2)    # => [1, 2]  
  
# Count number of elements  
> nums.count()    # => 3  
  
# Merge elements with an associative function  
> nums.reduce(lambda x, y: x + y) # => 6  
  
# Write elements to a text file  
> nums.saveAsTextFile("hdfs://file.txt")
```

# Working with Key-Value Pairs

Spark's “distributed reduce” transformations operate on RDDs of key-value pairs

**Python:**     `pair = (a, b)`  
                  `pair[0] # => a`  
                  `pair[1] # => b`

**Scala:**        `val pair = (a, b)`  
                  `pair._1 // => a`  
                  `pair._2 // => b`

**Java:**          `Tuple2 pair = new Tuple2(a, b);`  
                  `pair._1 // => a`  
                  `pair._2 // => b`



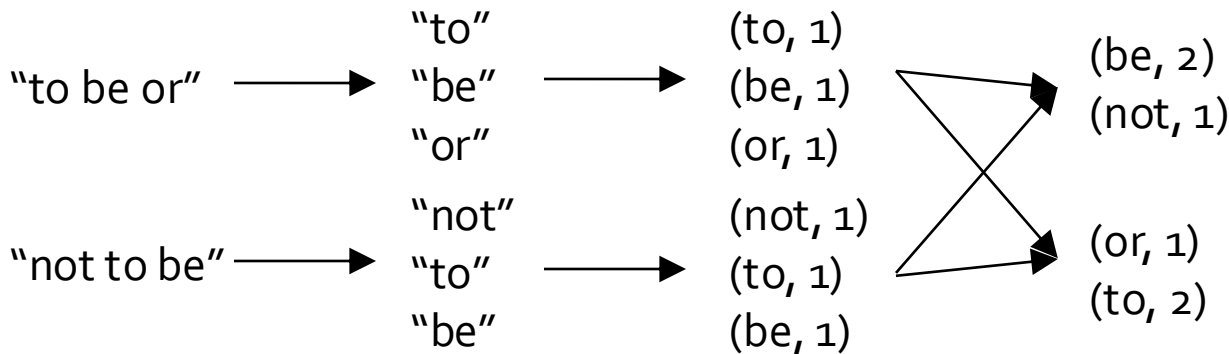
# Some Key-Value Operations

```
> pets = sc.parallelize(  
  [("cat", 1), ("dog", 1), ("cat", 2)])  
> pets.reduceByKey(lambda x, y: x + y)  
      # => {(cat, 3), (dog, 1)}  
> pets.groupByKey() # => {(cat, [1, 2]), (dog, [1])}  
> pets.sortByKey()  # => {(cat, 1), (cat, 2), (dog, 1)}
```

`reduceByKey` also automatically implements combiners on the map side

# Example: Word Count

```
> lines = sc.textFile("hamlet.txt")  
> counts = lines.flatMap(lambda line: line.split(" "))  
                    .map(lambda word: (word, 1))  
                    .reduceByKey(lambda x, y: x + y)
```



# Other Key-Value Operations

```
>visits = sc.parallelize([ ("index.html", "1.2.3.4"),  
                           ("about.html", "3.4.5.6"),  
                           ("index.html", "1.3.3.1") ])
```

```
>pageNames = sc.parallelize([ ("index.html", "Home"),  
                              ("about.html", "About") ])
```

```
>visits.join(pageNames)  
# ("index.html", ("1.2.3.4", "Home"))  
# ("index.html", ("1.3.3.1", "Home"))  
# ("about.html", ("3.4.5.6", "About"))
```

```
>visits.cogroup(pageNames)  
# ("index.html", ([ "1.2.3.4", "1.3.3.1"], [ "Home"] ))  
# ("about.html", ([ "3.4.5.6"], [ "About"] ))
```

# Setting the Level of Parallelism

All the pair RDD operations take an optional second parameter for number of tasks

```
> words.reduceByKey(lambda x, y: x + y, 5)
> words.groupByKey(5)
> visits.join(pageViews, 5)
```

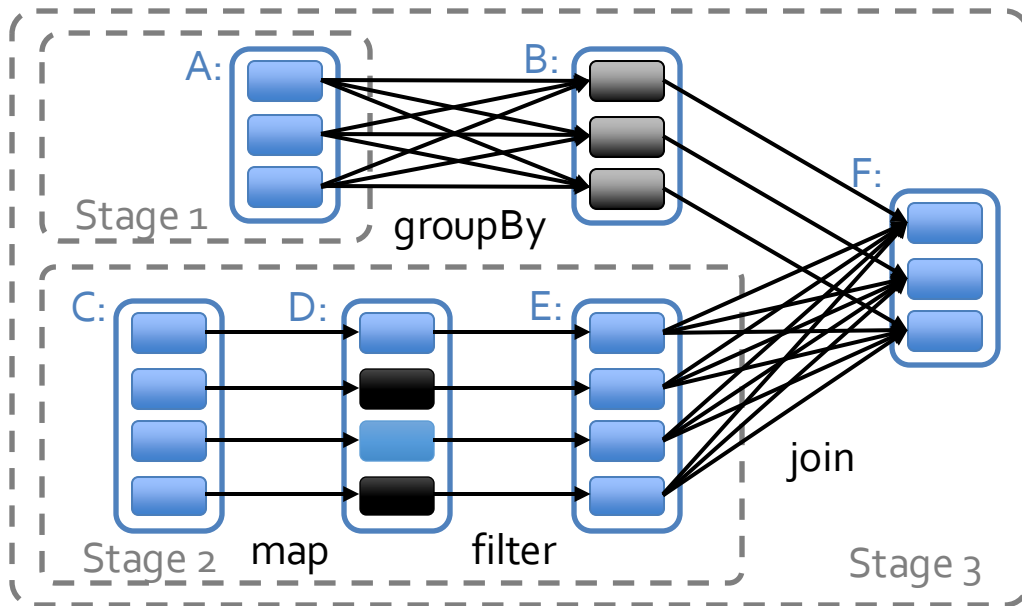
# Under The Hood: DAG Scheduler

General task  
graphs

Automatically  
pipelines functions

Data locality aware

Partitioning aware  
to avoid shuffles



= RDD



= cached partition

Marco Brambilla.

# More RDD Operators

map

filter

groupByKey

sort

union

join

leftOuterJoin

rightOuterJoin

reduce

count

fold

reduceByKey

groupByKey

cogroup

cross

zip

sample

take

first

partitionBy

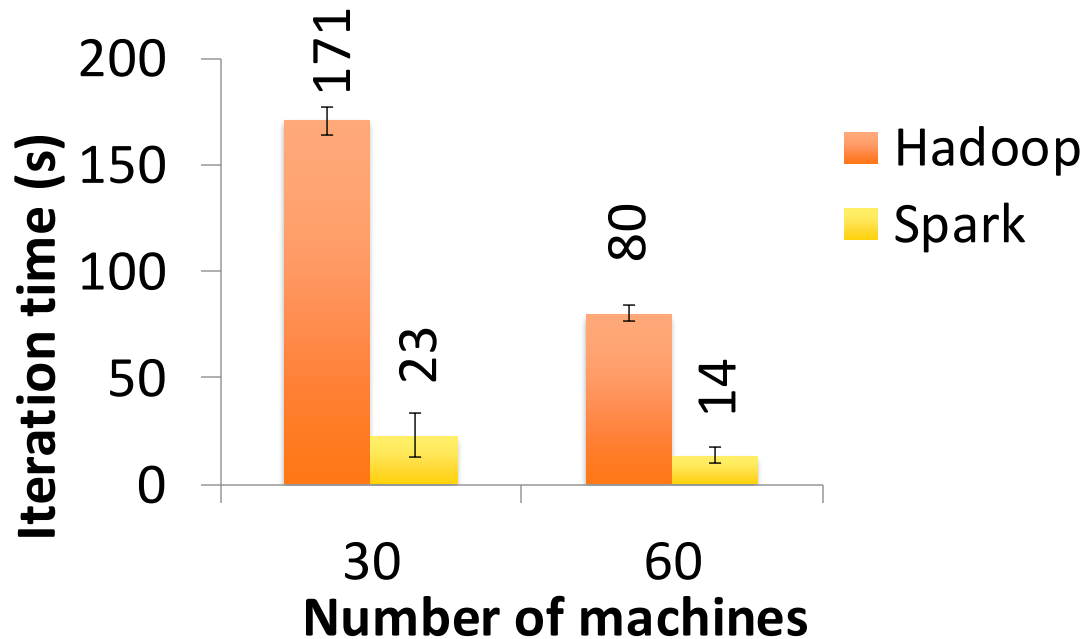
mapWith

pipe

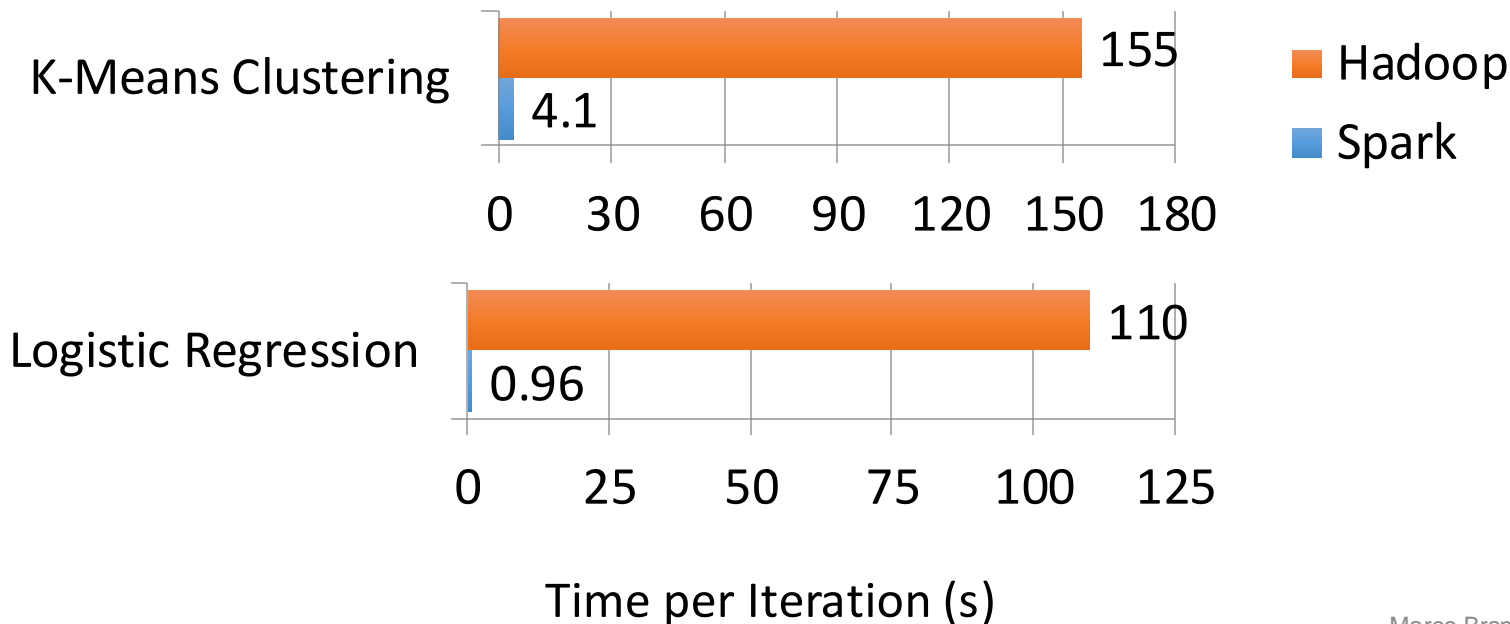
save

...

# PageRank Performance



# Other Iterative Algorithms





# On-Disk Sort Record:

## Time to sort 100TB

Record:  
Hadoop

2100 machines



72 minutes



Record:  
Spark

207 machines



23 minutes



Also sorted 1PB in 4 hours

Source: Daytona GraySort benchmark, [sortbenchmark.org](http://sortbenchmark.org)

Marco Brambilla

# Available File Formats

- Text / CSV
- JSON
- SequenceFile
  - binary key/value pair format
- Avro
- Parquet
- Data Frames
- ORC
  - optimized row columnar format

# AVRO

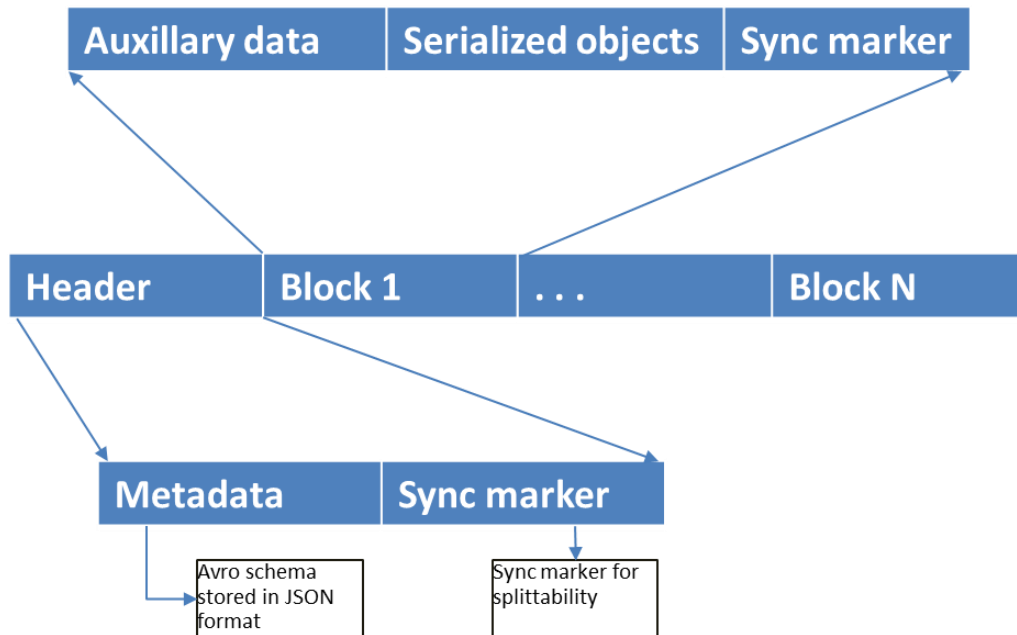
- Language neutral data serialization system
  - Write a file in python and read it in C
- AVRO data is described using language independent schema
- AVRO schemas are usually written in JSON and data is encoded in binary format
- Supports schema evolution
  - producers and consumers at different versions of schema
- Supports compression and are splittable

# Avro – File structure and example

Sample AVRO schema in JSON format

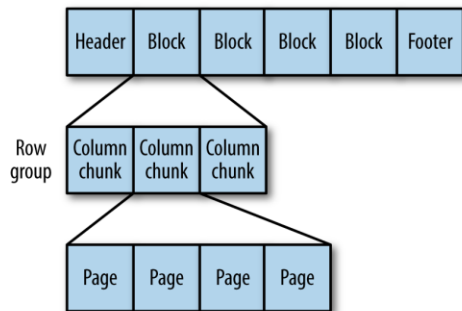
```
{
  "type" : "record",
  "name" : "tweets",
  "fields" : [ {
    "name" : "username",
    "type" : "string",
  }, {
    "name" : "tweet",
    "type" : "string",
  }, {
    "name" : "timestamp",
    "type" : "long",
  } ],
  "doc:" : "schema for storing tweets"
}
```

Avro file structure



# Parquet file structure & Configuration

Internal structure of parquet file



Configurable parquet parameters

Property name	Default value	Description
parquet.block.size	128 MB	The size in bytes of a block (row group) .
parquet.page.size	1MB	The size in bytes of a page .
parquet.dictionary.page.size	1MB	The maximum allowed size in bytes of a dictionary before falling back to plain encoding for a page.
parquet.enable.dictionary	true	Whether to use dictionary encoding.
parquet.compression	UNCOMPRESSED	The type of compression: UNCOMPRESSED, SNAPPY, GZIP & LZO

In summation, Parquet is state-of-the-art, open-source columnar format the supports *most* of processing frameworks and is optimized for high compression and high scan efficiency

# DataFrame

*noun* – [dey-tuh-freym]

1. A distributed collection of rows organized into named columns.
2. An abstraction for selecting, filtering, aggregating and plotting structured data (*cf. R, Pandas*).
3. Archaic: Previously SchemaRDD (*cf. Spark < 1.3*).

# DataFrame

```
ctx = new HiveContext()  
users = ctx.table("users")  
young = users.where(users("age") < 21)  
println(young.count())
```

- A distributed collection of rows with the same schema (RDDs suffer from type erasure)
- Can be constructed from external data sources or RDDs into essentially an RDD of Row objects (SchemaRDDs as of Spark < 1.3)
- Supports relational operators (e.g. *where*, *groupby*) as well as Spark operations.
- Evaluated lazily → unmaterialized *logical* plan

# Data Model

- Nested data model
- Supports
  - primitive SQL types (boolean, integer, double, decimal, string, data, timestamp)
  - complex types (structs, arrays, maps, and unions)
  - also user defined types.
- First class support for complex data types



# DataFrame Operations

- Relational operations (select, where, join, groupBy) via a DSL
- Operators take *expression* objects
- Operators build up an abstract syntax tree (AST), which is then optimized by *Catalyst*.

```
employees
  .join(dept, employees("deptId") === dept("id"))
  .where(employees("gender") === "female")
  .groupBy(dept("id"), dept("name"))
  .agg(count("name"))
```

- Alternatively, register as temp SQL table and perform traditional SQL query strings

```
users.where(users("age") < 21)
  .registerTempTable("young")
ctx.sql("SELECT count(*), avg(age) FROM young")
```

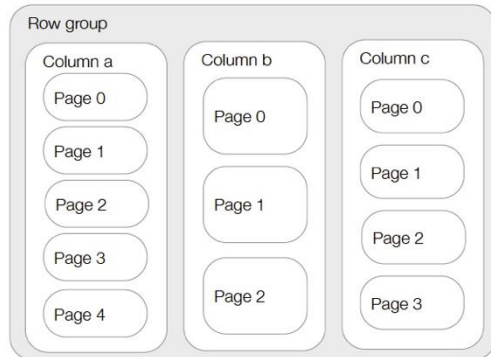
# Parquet

Apache Parquet is a **columnar storage** format available to any project in the Hadoop/Spark ecosystem, regardless of the choice of data processing framework, data model or programming language.



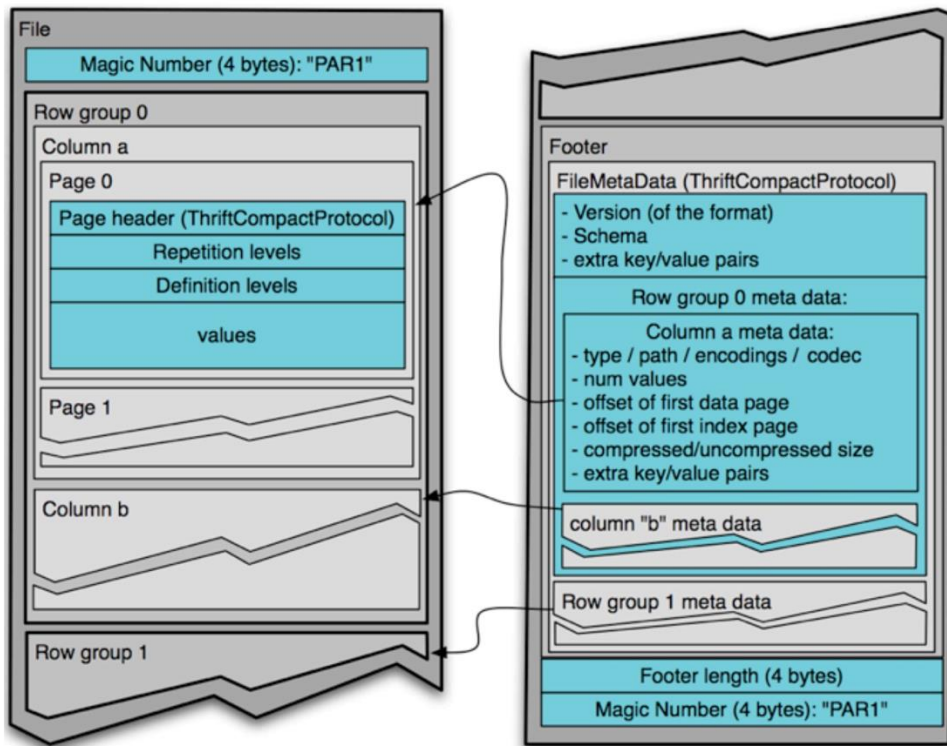
# Parquet Format

- **Row group:** A group of rows in columnar format.
  - Max size buffered in memory while writing.
  - One (or more) per split while reading.
  - roughly:  $50\text{MB} < \text{row group} < 1\text{ GB}$
- **Column chunk:** The data for one column in a row group.
  - Column chunks can be read independently for efficient scans.
- **Page:** Unit of access in a column chunk.
  - Should be big enough for compression to be efficient.
  - Minimum size to read to access a single record (when index pages are available).
  - roughly:  $8\text{KB} < \text{page} < 1\text{MB}$



Row group

# Parquet Format – details



- **Layout:**

Row groups in columnar format. A footer contains column chunks offset and schema.

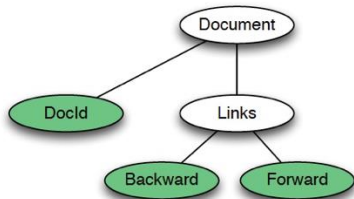
- **Language independent:**

Well defined format. Hadoop

# Nested Record Shredding & Assembly

Schema:

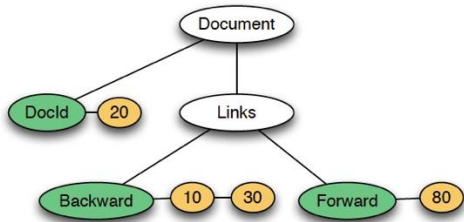
```
message Document {  
  required int64 DocId;  
  optional group Links {  
    repeated int64 Backward;  
    repeated int64 Forward;  
  }  
}
```



Columns	Max rep. level	Max def. level
DocId	0	0
Links.Backward	1	2
Links.Forward	1	2

Record:

```
DocId: 20  
Links  
  Backward: 10  
  Backward: 30  
  Forward: 80
```

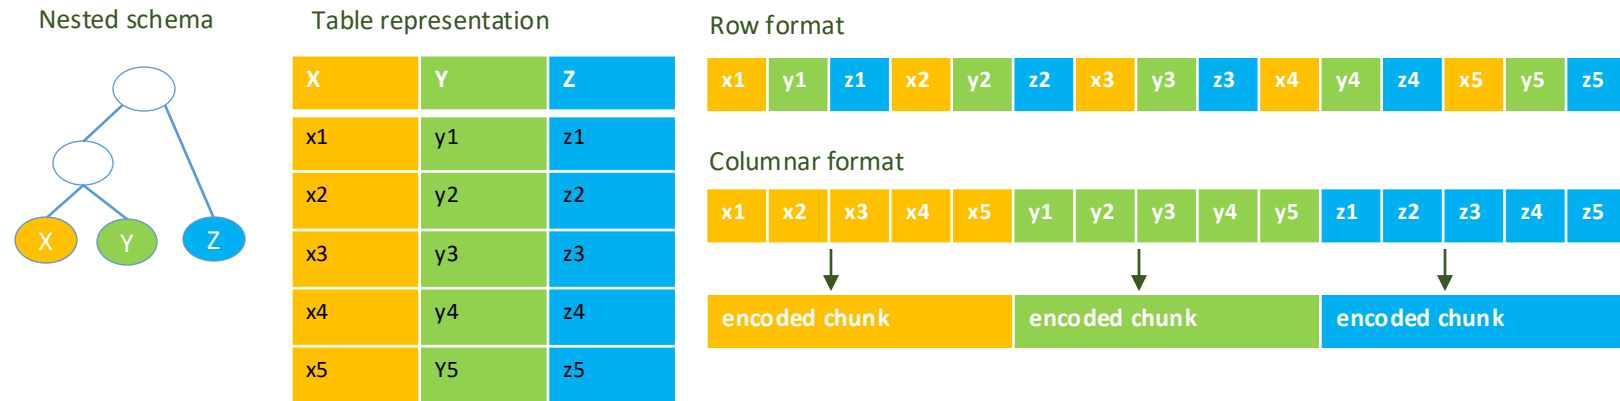


Column	Value	R	D
DocId	20	0	0
Links.Backward	10	0	2
Links.Backward	30	1	2
Links.Forward	80	0	2

# Parquet

columnar storage format

key strength is to store nested data in truly columnar format using definition and repetition levels<sup>1</sup>



(1) Dremel made simple with parquet - <https://blog.twitter.com/2013/dremel-made-simple-with-parquet>

# Optimizations – CPU and I/O

Statistics for filtering and query optimization

projection push down

X	Y	Z
x1	y1	z1
x2	y2	z2
x3	y3	z3
x4	y4	z4
x5	y5	z5

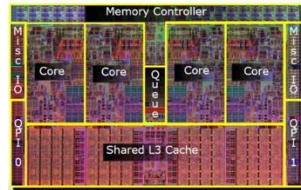
predicate push down

X	Y	Z
x1	y1	z1
x2	y2	z2
x3	y3	z3
x4	y4	z4
x5	y5	z5

read only the data you need

X	Y	Z
x1	y1	z1
x2	y2	z2
x3	y3	z3
x4	y4	z4
x5	y5	z5

Minimizes CPU cache misses



cache misses costs cpu cycles



# Encoding

## -Delta Encoding:

- E.g timestamp can be encoded by storing first value and the delta between subsequent values which tend to be small due to temporal validity

## -Prefix Encoding:

- delta encoding for strings

## -Dictionary Encoding:

- Small set of values, e.g post code, ip addresses etc

## -Run Length Encoding:

- repeating data



# Read Less Data

## Columnar organization

- Encoding: make the data smaller
- Column projection: read only the columns you need

## Row group filtering

- Use footer stats to eliminate row groups
- Use dictionary pages to eliminate row groups

## Page filtering

- Use page stats to eliminate pages

# Encoding

## Bit packing:

- Small integers encoded in the minimum bits required
- Useful for repetition level, definition levels and dictionary keys

## Run Length Encoding:

- Used in combination with bit packing
- Cheap compression
- Works well for definition level of sparse columns.

## Dictionary encoding:

- Useful for columns with few (  $< 50,000$  ) distinct values
- When applicable, compresses better and faster than heavyweight algorithms (gzip, lzo, snappy)

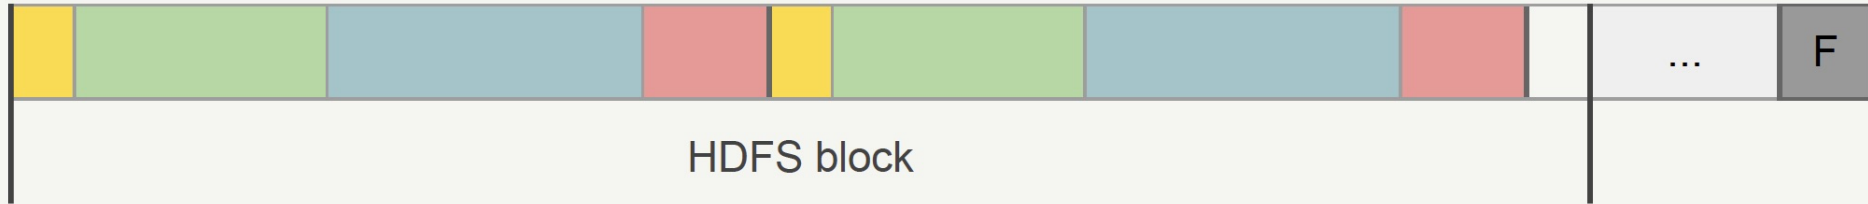
**Extensible:** Defining new encodings is supported by the format

# Parquet 2.0

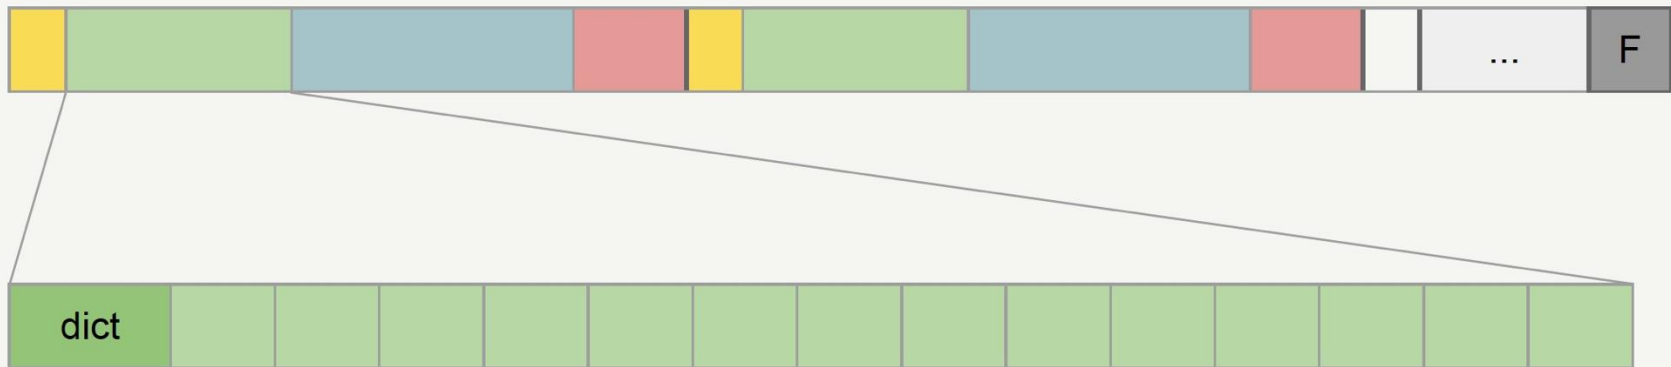
- More encodings: compact storage **without heavyweight compression**
- Delta encodings: for integers, strings and sorted dictionaries.
- Improved encoding for strings and boolean
- Statistics: to be used by query planners and predicate pushdown.
- New page format: to facilitate skipping ahead at a more granular level.

# .. On HDFS: Row Groups

A	B	C	D
a1	b1	c1	d1
...	...	...	...
aN	bN	cN	dN
...	...	...	...



# Column Chunks and Pages



Dictionary is a compact list of all the values.

- Search term missing? Skip the row group
- Like a bloom filter without false positives



is about more than SQL.

# Challenges and Solutions

## Challenges

- Perform ETL to and from various (semi- or unstructured) data sources
- Perform advanced analytics (e.g. machine learning, graph processing) that are hard to express in relational systems.

## Solutions

- A *DataFrame* API that can perform relational operations on both external data sources and Spark's built-in RDDs.
- A highly extensible optimizer, *Catalyst*, that uses features of Scala to add composable rule, control code gen., and define extensions.

# **Spark SQL** : Declarative BigData Processing

Let Developers Create and Run Spark Programs Faster:

- Write less code

- Read less data

- Let the optimizer do the hard work



# Write Less Code: Compute an Average



```
private IntWritable one =
    new IntWritable(1)
private IntWritable output =
    new IntWritable()
protected void map(
    LongWritable key,
    Text value,
    Context context) {
    String[] fields = value.split("\t")
    output.set(Integer.parseInt(fields[1]))
    context.write(one, output)
}

IntWritable one = new IntWritable(1)
DoubleWritable average = new DoubleWritable()

protected void reduce(
    IntWritable key,
    Iterable<IntWritable> values,
    Context context) {
    int sum = 0
    int count = 0
    for(IntWritable value : values) {
        sum += value.get()
        count++
    }
    average.set(sum / (double) count)
    context.write(key, average)
}
```



```
data = sc.textFile(...).split("\t")
data.map(lambda x: (x[0], [x[1], 1])) \
    .reduceByKey(lambda x, y: [x[0] + y[0], x[1] + y[1]]) \
    .map(lambda x: [x[0], x[1][0] / x[1][1]]) \
    .collect()
```

# Write Less Code: Compute an Average

## Using RDDs

```
data = sc.textFile(...).split("\t")
data.map(lambda x: (x[0], [int(x[1]), 1])) \
    .reduceByKey(lambda x, y: [x[0] + y[0], x[1] + y[1]]) \
    .map(lambda x: [x[0], x[1][0] / x[1][1]]) \
    .collect()
```

## Using SQL

```
SELECT name, avg(age)
FROM people
GROUP BY name
```

## Using Pig

```
P = load '/people' as (name, name);
G = group P by name;
R = foreach G generate ... AVG(G.age);
```

## Using DataFrames

```
sqlCtx.table("people") \
    .groupBy("name") \
    .agg("name", avg("age")) \
    .collect()
```

# Seamlessly Integrated: RDDs

Internally, DataFrame execution is done with Spark RDDs making interoperation with outside sources and custom algorithms easy.

## External Input

```
def buildScan(  
  requiredColumns: Array[String],  
  filters: Array[Filter]):  
  RDD[Row]
```

## Custom Processing

```
queryResult.rdd.mapPartitions { iter =>  
  ... Your code here ...  
}
```

# Extensible Input & Output

Spark's Data Source API allows optimizations like column pruning and filter pushdown into custom data sources.

## Built-In



## External




# Seamlessly Integrated

Embedding in a full programming language makes UDFs trivial and allows composition using functions.

```
zipToCity = udf(lambda city: <custom logic here>)
```

```
def add_demographics(events):  
    u = sqlCtx.table("users")  
    events \  
        .join(u, events.user_id == u.user_id) \  
        .withColumn("city", zipToCity(df.zip))
```



Takes and  
returns a  
DataFrame

# About *Spark* SQL

## Spark SQL

Part of the core distribution since Spark 1.0  
(April 2014)

Runs SQL / HiveQL queries, optionally  
alongside or replacing existing Hive  
deployments



```
SELECT COUNT(*)  
FROM hiveTable  
WHERE hive_udf(data)
```

# About *Spark*SQL

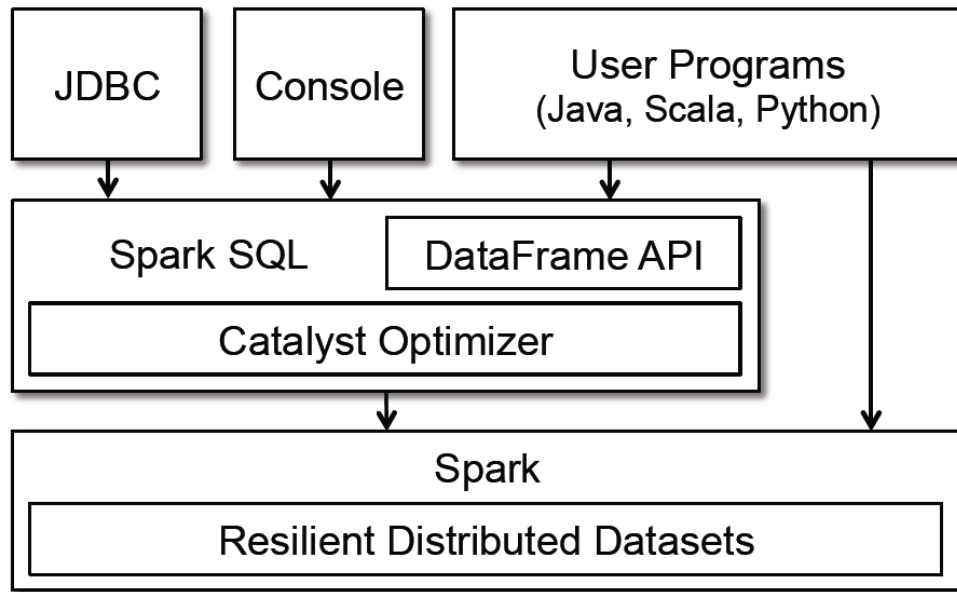
Originally called **SHARK**

Can only be used to query external data in Hive catalog → limited data sources

Can only be invoked via SQL string from Spark → error prone

Hive optimizer tailored for MapReduce → difficult to extend

# Programming Interface





# Advantages over Relational Query Languages

- Holistic optimization across functions composed in different languages.
- Control structures (e.g. *if*, *for*)
- Logical plan analyzed *eagerly* → identify code errors associated with data *schema* issues on the fly.

# Querying Native Datasets

- Infer column names and types directly from data objects (via reflection in Java and Scala and data sampling in Python, which is dynamically typed)

```
case class User(name: String, age: Int)
```

- Native objects accessed in-place to avoid expensive data format transformation.
- Benefits:
  - Run relational operations on existing Spark programs.
  - Combine RDDs with external structured data

Columnar storage with  
*hot* columns cached in  
memory

# User-Defined Functions (UDFs)

- Easy extension of limited operations supported.
- Allows inline registration of UDFs
  - Compare with Pig, which requires the UDF to be written in a Java package that's loaded into the Pig script.
- Can be defined on simple data types or entire tables.
- UDFs available to other interfaces after registration

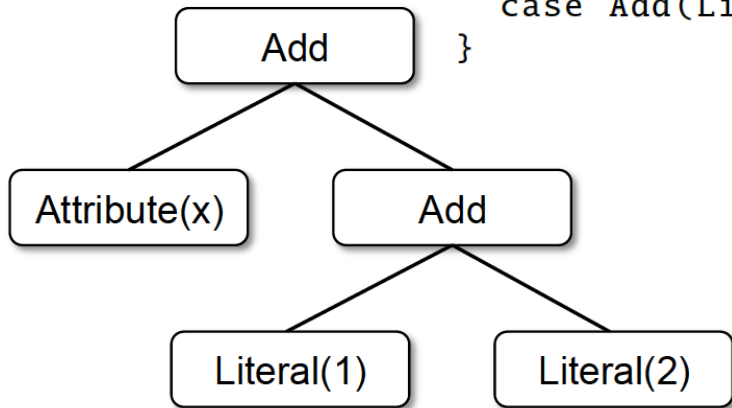
```
val model: LogisticRegressionModel = ...

ctx.udf.register("predict",
  (x: Float, y: Float) => model.predict(Vector(x, y)))

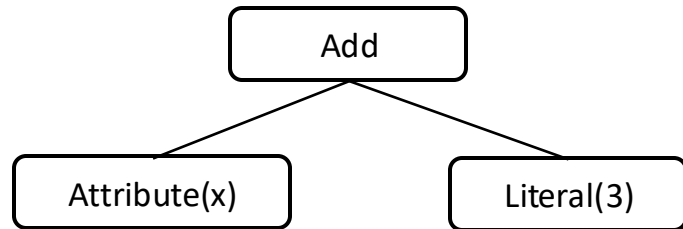
ctx.sql("SELECT predict(age, weight) FROM users")
```

# Query Optimization: Catalyst

```
tree.transform {  
  case Add(Literal(c1), Literal(c2)) => Literal(c1+c2)  
}
```



$x + (1 + 2)$

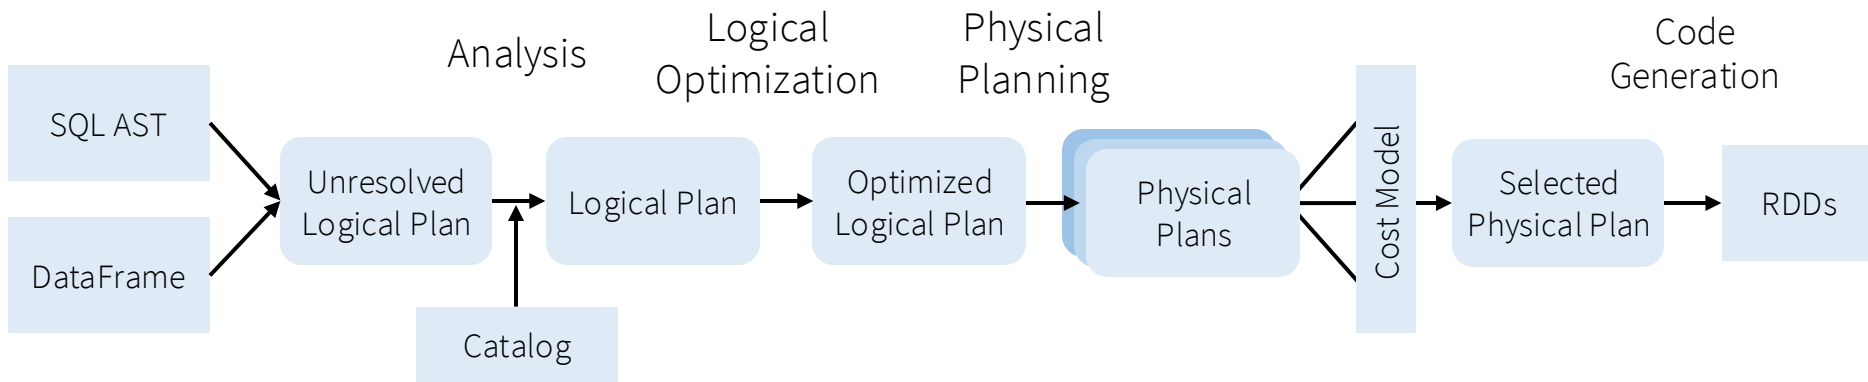


$x + 3$

# Catalyst Rules

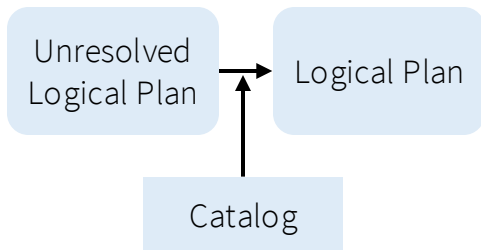
- *Pattern matching* functions that transform subtrees into specific structures.
  - Partial function*—skip over subtrees that do not match → no need to modify existing rules when adding new types of operators.
- Multiple patterns in the same *transform* call.
- May take multiple *batches* to reach a *fixed point*.
- *transform* can contain arbitrary Scala code.

# Plan Optimization & Execution



DataFrames and SQL share the same optimization/execution pipeline

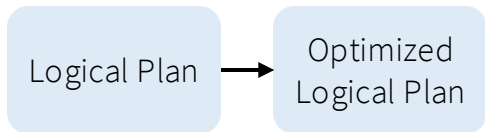
# Analysis



`SELECT col FROM sales`

- An attribute is *unresolved* if its type is not known or it's not matched to an input table.
- To resolve attributes:
  - Look up relations by name from the catalog.
  - Map named attributes to the input provided given operator's children.
  - UID for references to the same value
  - Propagate and coerce types through expressions (e.g.  $1 + col$ )

# Logical Optimization



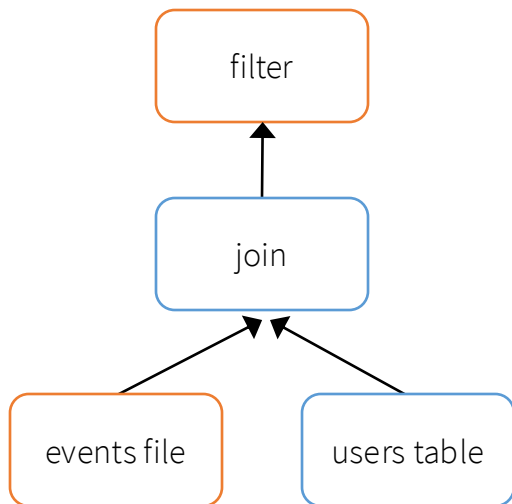
- Applies standard rule-based optimization (constant folding, predicate-pushdown, projection pruning, null propagation, boolean expression simplification, etc)
- 800LOC



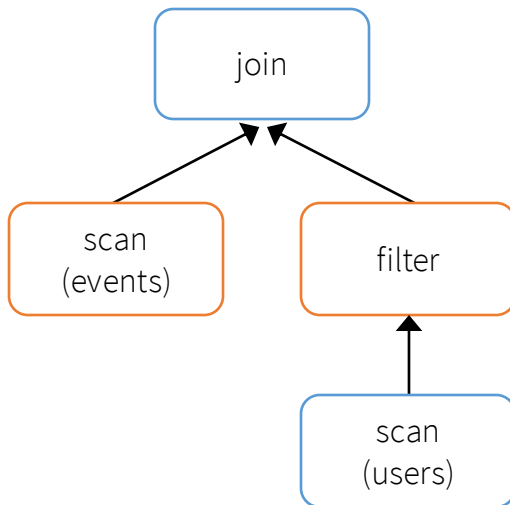
```
def add_demographics(events):
    u = sqlCtx.table("users")           # Load partitioned Hive table
    events \
        .join(u, events.user_id == u.user_id) \   # Join on user_id
        .withColumn("city", zipToCity(df.zip))      # Run udf to add city column

events = add_demographics(sqlCtx.load("/data/events", "parquet"))
training_data = events.where(events.city == "Melbourne").select(events.timestamp).collect()
```

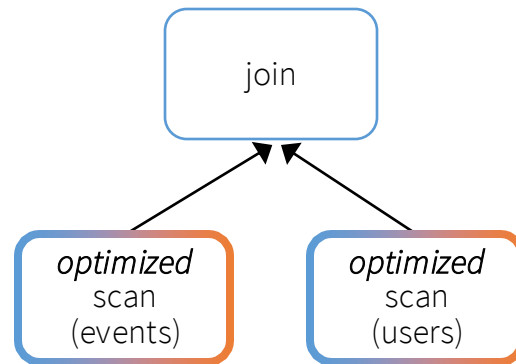
Logical Plan



## Physical Plan

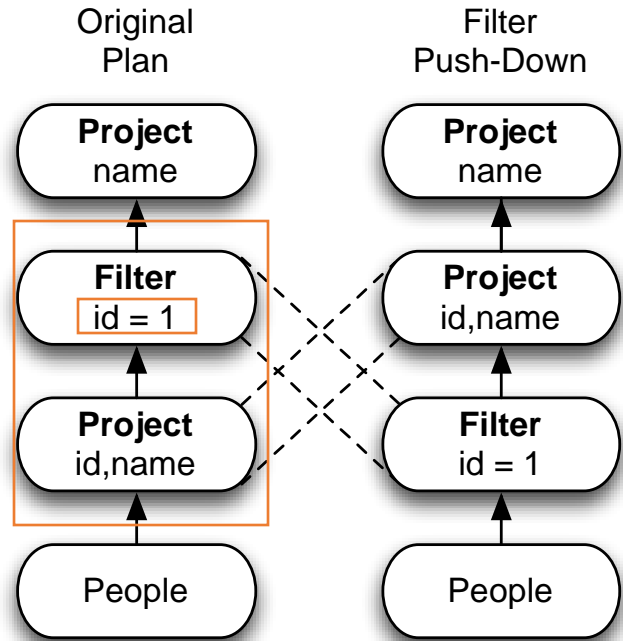


Physical Plan  
with Predicate Pushdown  
and Column Pruning



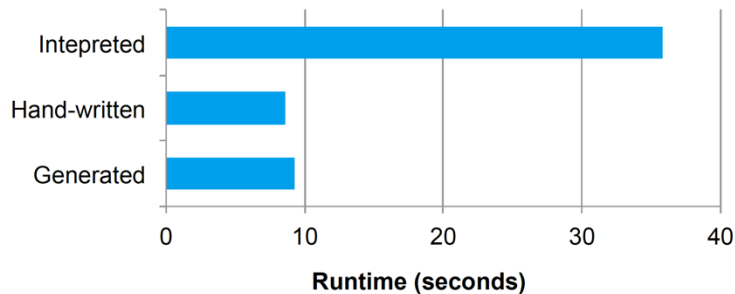
# An Example Catalyst Transformation

1. Find filters on top of projections.
2. Check that the filter can be evaluated without the result of the project.
3. If so, switch the operators.



# Code Generation

```
def compile(node: Node): AST = node match {  
  case Literal(value) => q"$value"  
  case Attribute(name) => q"row.get($name)"  
  case Add(left, right) =>  
    q"${compile(left)} + ${compile(right)}"  
}
```



**Figure 4:** A comparison of the performance evaluating the expression  $x+x+x$ , where  $x$  is an integer, 1 billion times.

- Relies on Scala's *quasiquotes* to simplify code gen.
- Catalyst transforms a SQL tree into an abstract syntax tree (AST) for Scala code to eval expr and generate code
- 700LOC

# Extensions

## Data Sources

- must implement a *createRelation* function that takes a set of key-value params and returns a *BaseRelation* object.
- E.g. CSV, Avro, Parquet, JDBC

## User-Defined Types (UDTs)

- Map user-defined types to structures composed of Catalyst's built-in types.

```
class PointUDT extends UserDefinedType[Point] {  
  def dataType = StructType(Seq( // Our native structure  
    StructField("x", DoubleType),  
    StructField("y", DoubleType)  
  ))  
  def serialize(p: Point) = Row(p.x, p.y)  
  def deserialize(r: Row) =  
    Point(r.getDouble(0), r.getDouble(1))  
}
```

# Advanced Analytics Features

## Schema Inference for Semistructured Data

### JSON

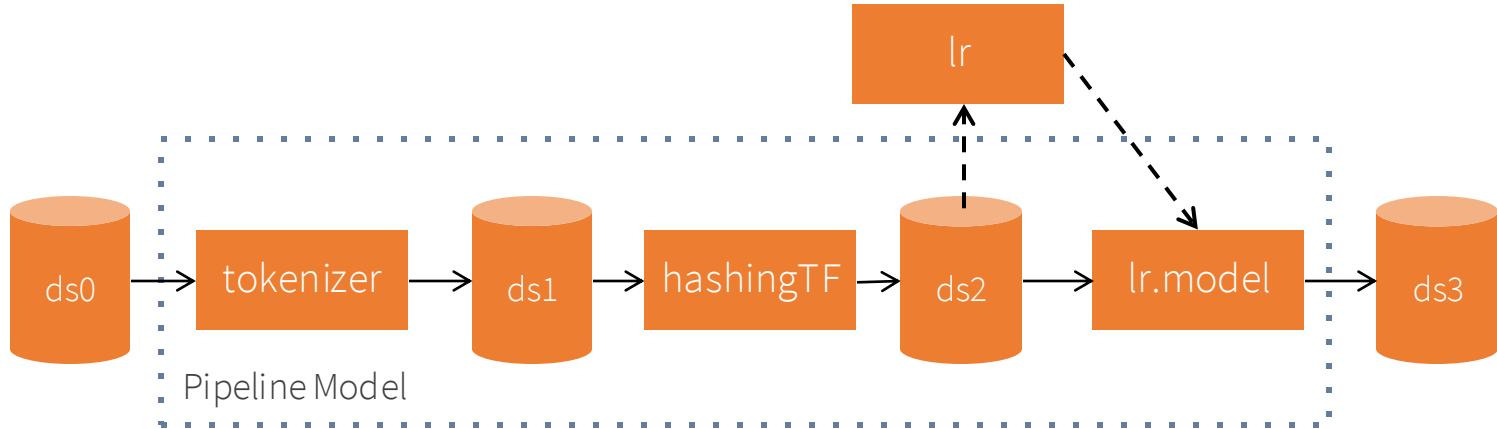
- Automatically infers schema from a set of records, in one pass or sample
- A tree of STRUCT types, each of which may contain atoms, arrays, or other STRUCTs.
- Find the most appropriate type for a field based on all data observed in that column. Determine array element types in the same way.
- Merge schemata of single records in one *reduce* operation.
- Same trick for Python typing

```
{  
  "text": "This is a tweet about #Spark",  
  "tags": ["#Spark"],  
  "loc": {"lat": 45.1, "long": 90}  
}  
  
{  
  "text": "This is another tweet",  
  "tags": [],  
  "loc": {"lat": 39, "long": 88.5}  
}  
  
{  
  "text": "A #tweet without #location",  
  "tags": ["#tweet", "#location"]  
}
```

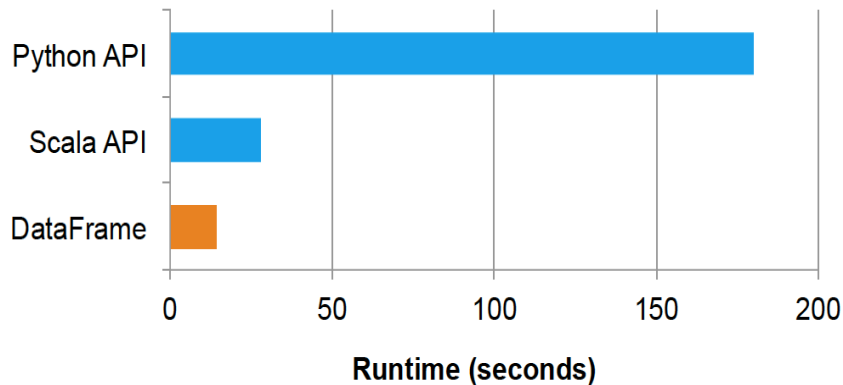
```
text STRING NOT NULL,  
tags ARRAY<STRING NOT NULL> NOT NULL,  
loc STRUCT<lat FLOAT NOT NULL, long FLOAT NOT NULL>
```

# Spark MLlib Pipelines

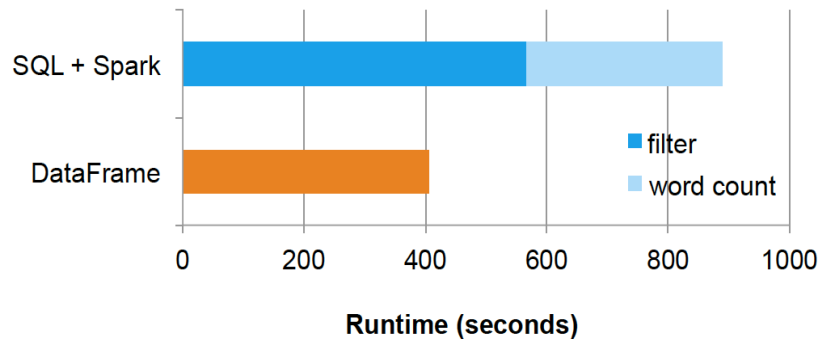
```
tokenizer = Tokenizer(inputCol="text", outputCol="words")  
hashingTF = HashingTF(inputCol="words", outputCol="features")  
lr = LogisticRegression(maxIter=10, regParam=0.01)  
pipeline = Pipeline(stages=[tokenizer, hashingTF, lr])  
  
df = sqlCtx.load("/path/to/data")  
model = pipeline.fit(df)
```



# About Performance Again..



**Figure 9: Performance of an aggregation written using the native Spark Python and Scala APIs versus the DataFrame API.**

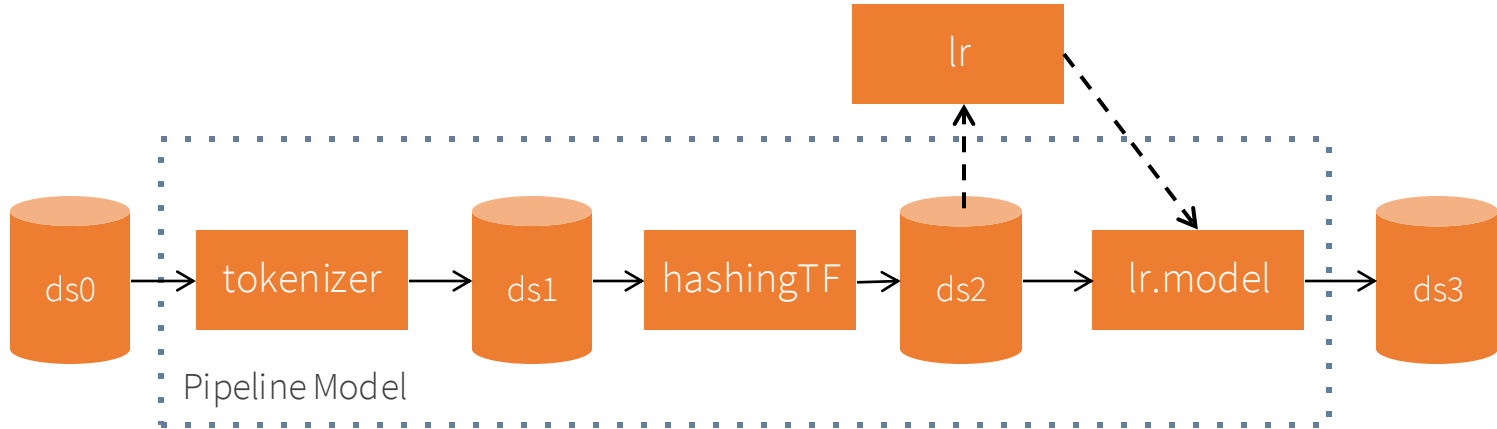


**Figure 10: Performance of a two-stage pipeline written as a separate Spark SQL query and Spark job (above) and an integrated DataFrame job (below).**

# Spark MLlib Pipelines

```
tokenizer = Tokenizer(inputCol="text", outputCol="words")
hashingTF = HashingTF(inputCol="words", outputCol="features")
lr = LogisticRegression(maxIter=10, regParam=0.01)
pipeline = Pipeline(stages=[tokenizer, hashingTF, lr])

df = sqlCtx.load("/path/to/data")
model = pipeline.fit(df)
```





**<http://spark.apache.org/>**



**POLITECNICO**  
MILANO 1863

SYSTEMS AND METHODS FOR BIG AND UNSTRUCTURED DATA

# Thanks

Marco Brambilla

marco.brambilla@polimi.it

 @marcobrambi